

$$1 \times 2^1 + 0 \times 2^0 = 2$$

There are only 10 types of students

- those that understand binary
- those that don't understand binary

- Can interpret decimal number 4705 as:
 $4 \times 10^3 + 7 \times 10^2 + 0 \times 10^1 + 5 \times 10^0$

- The *base* or *radix* is 10
 Digits 0 – 9

- Place values:

...	1000	100	10	1
...	10^3	10^2	10^1	10^0

- Write number as 4705_{10}

- Note use of subscript to denote base

1

2

Binary Representation

Computer uses base 2.

- In a similar way, can interpret binary number 1011 as:

$$1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 = 11_{10}$$

- The *base* or *radix* is 2

Digits 0 and 1

- Place values:

...	8	4	2	1
...	2^3	2^2	2^1	2^0

- Write number as 1011_2
 $(= 11_{10})$

Hexadecimal Representation

- Can interpret hexadecimal number 3AF1 as:

$$3 \times 16^3 + 10 \times 16^2 + 15 \times 16^1 + 1 \times 16^0$$

- The *base* or *radix* is 16

Digits 0,1,2,3,4,5,6,7,8,9,A,B,C,D,E,F

- Place values:

...	4096	256	16	1
...	16^3	16^2	16^1	16^0

- Write number as $3AF1_{16}$
 $(= 15089_{10})$

3

4

In hexadecimal, each digit represents 4 bits

$$0 \times 2^3 + 1 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 \\ 4 + 2 + 1 = 7$$

Symbol for hexadecimal (16)

0100	1000	1111	1010	1011	1100	1001	0111	
0x	4	8	F	A	B	C	9	7

In octal, each digit represents 3 bits

01	001	000	111	110	101	011	110	010	010	111	
0	1	1	0	7	6	5	3	6	2	2	7

In binary, each digit represents 1 bit

0b0100100011110101011110010010111

- Example: Convert 101111000101001_2 to Hex:
- Example: Convert 10111101011100_2 to Hex:

5

6

Hexadecimal to Binary

- Reverse the previous process
- Convert each hex digit into equivalent 4-bit binary representation
- Example: Convert $AD5_{16}$ to Binary:

Bits in Bytes in Words

Values that we normally treat as atomic can be viewed as bits, e.g.

- char = 1 byte = 8 bits (a is 01100001)
- short = 2 bytes = 16 bits (42 is 0000000000101010)
- int = 4 bytes = 32 bits (42 is 00000000...0000101010)
- double = 8 bytes = 64 bits

The above are common sizes and don't apply on all hardware
e.g. `sizeof(int)` might be 2, 4 or 8.

C provides a set of operators that act bit-by-bit on pairs of bytes.

E.g. `(10101010 & 11110000) == 10100000` (bitwise AND)

C bitwise operators: & | ^ ~ << >>

The & operator

- takes two values (1,2,4,8 bytes), treats as sequence of bits
- performs logical AND on each corresponding pair of bits
- result contains same number of bits as inputs

Example:

00100111	AND 0 1	Only 2 "1"
& 11000011	---- -----	
-----	0 0 0	will result "1"
00100011	1 0 1	

Used for e.g. checking whether a bit is set

The obvious way to check for odd numbers in C

```
int isOdd(int n) {
    return n % 2 == 1;
}
```

We can use & to achieve the same thing:

```
int isOdd(int n) {
    return n & 1;
}
```

No special/extra benefits
to use this.
just need to aware this

Because

```
107 01101010
108 01101100
109 01101101
110 01101110
111 01101111
112 01110000
113 01110001
114 01110010
115 01110011
116 01110100
117 01110101
118 01110110
119 01110111
120 01111000
121 01111001
122 01111010
123 01111011
124 01111100
125 01111101
126 01111110
127 01111111
```

Even

Odd

0 & 1 => 0 => Not odd.

1 & 1 => 1 => Odd.

Bitwise OR

The | operator

- takes two values (1,2,4,8 bytes), treats as sequence of bits
- performs logical OR on each corresponding pair of bits
- result contains same number of bits as inputs

Example:

00100111	OR 0 1	Any "1"
11000011	---- -----	
-----	0 0 1	will result "1"
11100111	1 1 1	

Used for e.g. ensuring that a bit is set

Bitwise NEG

The ~ operator

- takes a single value (1,2,4,8 bytes), treats as sequence of bits
- performs logical negation of each bit
- result contains same number of bits as input

Example:

~ 00100111	NEG 0 1	<u>Opposite</u> .
-----	---- -----	
	11011000	1 0

Used for e.g. creating useful bit patterns

- everything is ultimately a string of bits
- e.g. `unsigned char` = 8-bit value
- e.g. literal bit-string `0b01110001`
- e.g. literal hexadecimal `0x71`
- `&` = bitwise AND
- `|` = bitwise OR
- `~` = bitwise NEG

The `^` operator

- takes two values (1,2,4,8 bytes), treats as sequence of bits
- performs logical XOR on each corresponding pair of bits
- result contains same number of bits as inputs

Example:

<code>00100111</code>	<code>XOR 0 1</code>	<i>Same number $\Rightarrow 0$</i>
<code>^ 11100011</code>	<code>----- -----</code>	<i>Different number $\Rightarrow 1$</i>
<code>-----</code>	<code>0 0 1</code>	
<code>11000100</code>	<code>1 1 0</code>	

Used in e.g. generating hashes, graphic operation, cryptography

5

6

Left Shift

The `<<` operator

- takes a single value (1,2,4,8 bytes), treats as sequence of bits
- and a small positive integer $\{x\}$
- moves (shifts) each bit $\{x\}$ positions to the left
- left-end bit vanishes; right-end bit replaced by zero
- result contains same number of bits as input

Example:

<code>00100111 << 2</code>	<code>00100111 << 8</code>
<code>-----</code>	<code>-----</code>
<code>10011100</code>	<code>00000000</code>

Right Shift

The `>>` operator

- takes a single value (1,2,4,8 bytes), treats as sequence of bits
- and a small positive integer $\{x\}$
- moves (shifts) each bit $\{x\}$ positions to the right
- right-end bit vanishes; left-end bit replaced by zero**
- result contains same number of bits as input

Example:

<code>00100111 >> 2</code>	<code>00100111 >> 8</code>
<code>-----</code>	<code>-----</code>
<code>00001001</code>	<code>00000000</code>

Beware: shifts involving negative values are not portable (implementation defined) - use unsigned values to be safe/portable.

7

8

Given the following variable declarations:

```
// a signed 8-bit value  
unsigned char x = 0x55;  
unsigned char y = 0xAA;
```

What is the value of each of the following expressions:

- $(x \& y)$ $(x ^ y)$
- $(x \ll 1)$ $(y \ll 1)$
- $(x \gg 1)$ $(y \gg 1)$

Assuming 8-bit quantities and writing answers as 8-bit bit-strings:

What are the values of the following:

- 25, 65, ~0, ~~1, 0xFF, ~0xFF
- (01010101 & 10101010), (01010101 | 10101010)
- $(x \& \sim x)$, $(x | \sim x)$

How can we achieve each of the following:

- ensure that the 3rd bit from the RHS is set to 1
- ensure that the 3rd bit from the RHS is set to 0

- C has three floating point types
 - float ... typically 32-bit (lower precision, narrower range)
 - double ... typically 64-bit (higher precision, wider range)
 - long double ... typically 128-bits (but maybe only 80 bits used)
- Floating point constants, e.g.: 3.14159 1.0e-9 are double
- Reminder: division of 2 ints in C yields an int.
 - but division of double and int in C yields a double.

```
double d = 4/7.0;
// prints in decimal with (default) 6 decimal places
printf("%lf\n", d);           // prints 0.571429
// prints in scientific notation
printf("%le\n", d);          // prints 5.714286e-01
// picks best of decimal and scientific notation
printf("%lg\n", d);          // prints 0.571429
// prints in decimal with 9 decimal places
printf("%.9lf\n", d);         // prints 0.571428571
// prints in decimal with 1 decimal place and field width of 5
printf("%10.1lf\n", d);       // prints      0.6
```

source code for float_output.c

1

2

- if we represent floating point numbers with a fixed small number of bits
 - there are only a finite number of bit patterns
 - can only represent a finite subset of reals
- almost all real values will have no exact representation
- value of arithmetic operations may be real with no exactly representation
- we must use closest value which can be exactly represented
- this approximation introduces an error into our calculations
- often does not matter
- sometime can be disasterous

$$\begin{array}{r} 1 \times 2^3 + 1 \times 2^1 + 0 \times 2^0 \\ \hline 6 \end{array} \quad \begin{array}{r} 1 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3} \\ \hline . \cdot \frac{1}{2} + \frac{1}{4} + \frac{1}{8} = 0.625_{10} \end{array}$$

- can have fractional numbers in other bases, e.g.: 110.101₂ == 6.625₁₀
- could represent fractional numbers similarly to integers by assuming decimal point is in **fixed** position
- for example with 32 bits:
 - 16 bits could be used for integer part
 - 16 bits could be used for the fraction
 - equivalent to storing values as integers after multiplying (*scaling*) by 2¹⁶
 - major limitation is only small range of values can be represented
 - minimum 2₋₁₆ ≈ 0.000015
 - maximum 2₁₅ ≈ 32768
- usable for some problems, but not ideal
- used on small embedded processors without silicon floating point

```

float f;
double d;
long double l;
printf("float      %2lu bytes min=%-12g  max=%g\n", sizeof f, F1)
printf("double     %2lu bytes min=%-12g  max=%g\n", sizeof d, DB1)
printf("long double %2lu bytes min=%-12Lg  max=%Lg\n", sizeof l, L1)

source code for floating_types.c

$ ./floating_types
float      4 bytes min=1.17549e-38  max=3.40282e+38
double     8 bytes min=2.22507e-308 max=1.79769e+308
long double 16 bytes min=3.3621e-4932 max=1.18973e+4932

```

In the lecture, he focus in float & single precision (binary32).
because double can be understood once float is understood.

- C floats almost always IEEE 754 single precision (binary32)
- C double almost always IEEE 754 double precision (binary64)
- C long double might be IEEE 754 (binary128)
- IEEE 754 representation has 3 parts: *sign*, *fraction* and *exponent*
- numbers have form $\text{sign} \frac{\text{fraction}}{2^{\text{exponent}}}$, where *sign* is +/-
- **fraction** always has 1 digit before decimal point (**normalized**)
 - as a consequence only 1 representation for any value
- **exponent** is stored as positive number by adding constant value (**bias**)
- numbers close to zero have higher precision (more accurate)

5

6

Floating Point Numbers

Example of normalising the fraction part in binary:

- 1010.1011 is normalized as 1.0101011×2^{011}
- $1010.1011 = 10 + 11/16 = 10.6875$
- $1.0101011 \times 2^{011} = (1 + 43/128) \times 2^3 = 1.3359375 \times 8 = 10.6875$

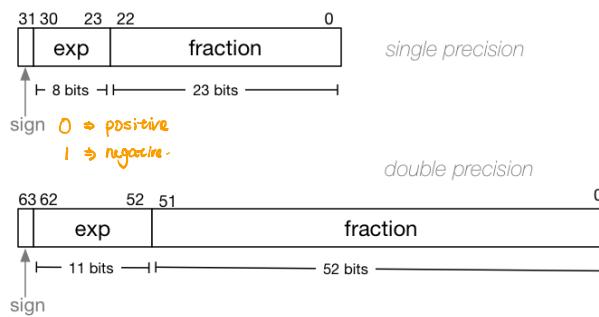
The normalised fraction part always has 1 before the decimal point. Assume to be there

Example of determining the exponent in binary:

- if exponent is 8-bits, then the bias = $2^{8-1} - 1 = 127$
- valid bit patterns for exponent 00000001 .. 11111110
- correspond to B exponent values -126 .. 127

Floating Point Numbers

Internal structure of floating point values



7

8

IEEE-754 Single Precision example: 0.15625

0.15625 is represented in IEEE-754 single-precision by these bits:
00111100010000000000000000000000
sign | exponent | fraction
0 | 0111100 | 010000000000000000000000
sign bit = 0
sign = +
raw exponent = 0111100 binary
= 124 decimal
actual exponent = 124 - exponent_bias
= 124 - 127
= -3
number = +1.0100000000000000000000000000000 binary * 2**-3
= 1.25 decimal * 2**-3
= 1.25 * 0.125
= 0.15625

source code for explain_float_representation.c

IEEE-754 Single Precision example: -0.125

\$./explain_float_representation -0.125
-0.125 is represented as a float (IEEE-754 single-precision) by the
10111100000000000000000000000000
sign | exponent | fraction
1 | 0111100 | 000000000000000000000000
sign bit = 1
sign = -
raw exponent = 0111100 binary
= 124 decimal
actual exponent = 124 - exponent_bias
= 124 - 127
= -3
number = -1.0000000000000000000000000000000 binary * 2**-3
= -1 decimal * 2**-3
= -1 * 0.125
= -0.125

9

10

IEEE-754 Single Precision example: 150.75

\$./explain_float_representation 150.75
150.75 is represented in IEEE-754 single-precision by these bits:
0100001100101101000000000000000
sign | exponent | fraction
0 | 10000110 | 001011011000000000000000
sign bit = 0
sign = +
raw exponent = 10000110 binary
= 134 decimal
actual exponent = 134 - exponent_bias
= 134 - 127
= 7
number = +1.001011011000000000000000 binary * 2**7
= 1.17773 decimal * 2**7
= 1.17773 * 128
= 150.75

IEEE-754 Single Precision example: -96.125

\$./explain_float_representation -96.125
-96.125 is represented in IEEE-754 single-precision by these bits:
1100001011000000100000000000000
sign | exponent | fraction
1 | 10000101 | 100000010000000000000000
sign bit = 1
sign = -
raw exponent = 10000101 binary
= 133 decimal
actual exponent = 133 - exponent_bias
= 133 - 127
= 6
number = -1.100000010000000000000000 binary * 2**6
= -1.50195 decimal * 2**6
= -1.50195 * 64
= -96.125

11

12

```
$ ./explain_float_representation 00111101110011001100110011001101
sign bit = 0
sign = +
raw exponent      = 01111011 binary
                  = 123 decimal
actual exponent = 123 - exponent_bias
                  = 123 - 127
                  = -4
number = +1.10011001100110011001101 binary * 2**-4
       = 1.6 decimal * 2**-4
       = 1.6 * 0.0625
       = 0.1
```

- C (IEEE-754) has a representation for +/- infinity
- propagates sensibly through calculations

```
double x = 1.0/0.0;
printf("%lf\n", x); //prints inf
printf("%lf\n", -x); //prints -inf
printf("%lf\n", x - 1); // prints inf
printf("%lf\n", 2 * atan(x)); // prints 3.141593
printf("%d\n", 42 < x); // prints 1 (true)
printf("%d\n", x == INFINITY); // prints 1 (true)
```

source code for infinity.c

13

14

nan.c - handling errors robustly

IEEE-754 Single Precision example: inf

- C (IEEE-754) has a representation for invalid results:
 - NaN (not a number)
- ensures errors propagate sensibly through calculations

```
double x = 0.0/0.0;
printf("%lf\n", x); //prints nan
printf("%lf\n", x - 1); // prints nan
printf("%d\n", x == x); // prints 0 (false)
printf("%d\n", isnan(x)); // prints 1 (true)

source code for nan.c
```

```
$ ./explain_float_representation inf
inf is represented in IEEE-754 single-precision by these bits:
01111111000000000000000000000000
sign | exponent | fraction
0 | 1111111 | 00000000000000000000000000000000
sign bit = 0
sign = +
raw exponent      = 11111111 binary
                  = 255 decimal
number = +inf
```

15

16

```
$ ./explain_float_representation 01111111100000000000000000000000
sign bit = 0
sign = +
raw exponent = 11111111 binary
              = 255 decimal
number = NaN

source code for explain_float_representation.c
```

exp.

```
double a, b;
a = 0.1;
b = 1 - (a + a + a + a + a + a + a + a + a + a);
if (b != 0) { // better would be fabs(b) > 0.000001
    printf("1 != 0.1+0.1+0.1+0.1+0.1+0.1+0.1+0.1\n");
}
printf("b = %g\n", b); // prints 1.11022e-16

source code for double_imprecision.c
```

- do not use == and != with floating point values
- instead check if values are close

17

Consequences of most reals not having exact representations

```
double x = 0.00000001;
double y = (1 - cos(x)) / (x * x);
// correct answer y = -0.5
// prints y = 0.917540
printf("y = %f\n", y);
// division of similar approximate value
// produces large error
// sometimes called catastrophic cancellation
printf("%g\n", 1 - cos(x)); // prints 1.11022e-16
printf("%g\n", x * x); // prints 1.21e-16

source code for double_catastrophe.c
```

Another reason not to use == with floating point values

```
if (d == d) {
    printf("d == d is true\n");
} else {
    // will be executed if d is a NaN
    printf("d == d is not true\n");
}
if (d == d + 1) {
    // may be executed if d is large
    // because closest possible representation for d + 1
    // is also closest possible representation for d
    printf("d == d + 1 is true\n");
} else {
    printf("d == d + 1 is false\n");
}

source code for double_not_always.c
```

19

18

20

Another reason not to use == with floating point values

```
$ gcc double_not_always.c -o double_not_always
$ ./double_not_always 42.3
d = 42.3
d == d is true
d == d + 1 is false
$ ./double_not_always 42000000000000000000
d = 4.2e+18
d == d is true
d == d + 1 is true
$ ./double_not_always NaN
d = nan
d == d is not true
d == d + 1 is false
```

because closest possible representation for $d + 1$ is also closest possible representation for d

source code for double_not_always.c

Exercise: Floating point → Decimal

Convert the following floating point numbers to decimal.

Assume that they are in IEEE 754 single-precision format.

0 10000000 11000000000000000000000000000000

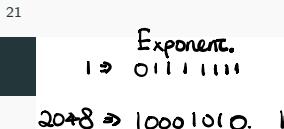
1 01111110 10000000000000000000000000000000

Consequences of most reals not having exact representations

```
double d = 9007199254740992;
// loop never terminates
while (d < 9007199254740999) {
    printf("%lf\n", d); // always prints 9007199254740992.000000
    // 9007199254740993 can not be represented as a double
    // closest double is 9007199254740992.0
    // so 9007199254740992.0 + 1 = 9007199254740992.0
    d = d + 1;
}
```

source code for double_disaster.c

- 9007199254740993 is $2^{53} + 1$
it is smallest integer which can not be represented exactly as a double
- The closest double to 9007199254740993 is 9007199254740992.0
- aside: 9007199254740993 can not be represented by a int32_t
it can be represented by int64_t



21

22

23

Useful to know assembly language because ...

- sometimes you are *required* to use it (e.g. device handlers)
- improves your understanding of how compiled programs execute
 - very helpful when debugging
 - understand performance issues better
- performance tweaking (squeezing out last pico-second)
 - re-write that performance critical code in assembler

A typical modern CPU has

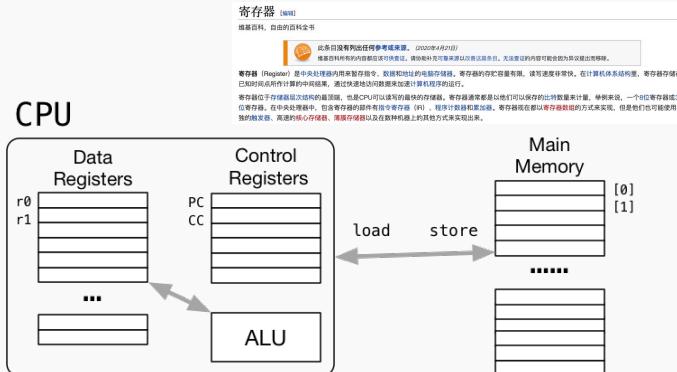
- a set of data registers
- a set of control registers (incl PC)
- an arithmetic-logic unit (ALU)
- access to memory (RAM)
- a set of simple instructions
 - transfer data between memory and registers
 - push values through the ALU to compute results
 - make tests and transfer control of execution

Different types of processors have different configurations of the above

1

2

CPU Architecture



Fetch-Execute Cycle

- typical CPU program execution pseudo-code:

```
uint32_t pc = STARTING_ADDRESS;
while (1) {
    uint32_t instruction = memory[pc];
    pc++; // move to next instr
    if (instruction == HALT) {
        break;
    } else {
        execute(instruction);
    }
}

▪ pc = Program Counter, a CPU register which tracks execution
  ▪ note some instructions modify pc (branches and jumps)
```

3

4

Executing an instruction involves

- determine what the *operator* is
- determine which *registers*, if any, are involved
- determine which *memory location*, if any, is involved
- carry out the operation with the relevant operands
- store result, if any, in appropriate register

Example instruction encodings (not from a real machine):

ADD	R1	R2	R3
— 8 bits —	— 8 bits —	— 8 bits —	— 8 bits —
LOAD	R4	0x10004	
— 8 bits —	— 8 bits —	— 16 bits —	

MIPS Instructions

MIPS has several classes of instructions:

- *load and store* ... transfer data between registers and memory
- *computational* ... perform arithmetic/logical operations
- *jump and branch* ... transfer control of program execution
- *coprocessor* ... standard interface to various co-processors
- *special* ... miscellaneous tasks (e.g. syscall)

And several *addressing modes* for each instruction

- between memory and register (direct, indirect)
- constant to register (immediate)
- register + register + destination register

MIPS is a well-known and simple architecture

- historically used everywhere from supercomputers to PlayStations, ...
- still popular in some embedded fields e.g. modems, TVs
- but being out-competed by arm (in phones, ...)

We consider the MIPS32 version of the MIPS family

- qtspim ... provides a GUI front-end, useful for debugging
- spim ... command-line based version, useful for testing
- xspim ... GUI front-end, useful for debugging, only in CSE labs

Executables and source: <http://spimsimulator.sourceforge.net/>

Source code for browsing under `/home/cs1521/spim`

5

MIPS Instructions

MIPS instructions are 32-bits long, and specify ...

- an operation (e.g. load, store, add, branch, ...)
- one or more operands (e.g. registers, memory addresses, constants)

Some possible instruction formats

OPCODE	R1	R2	R3	unused
— 6 bits —	— 5 bits —	— 5 bits —	— 5 bits —	— 11 bits —

OPCODE	R1	Memory Address or Constant Value
— 6 bits —	— 5 bits —	— 21 bits —

7

8

Assembly Language

- Instructions are simply bit patterns (32-bits on the MIPS)
- Could write machine code program just by specifying the bit-pattern e.g as a sequence of hex digits:

0x3c041001

0x34020004

0x0000000c

0x03e00008

- unreadable and difficult to maintain
- adding/removing instructions changes bit pattern for other instructions
- changing variable layout in memory, changes bit pattern for instructions
- solution: **assembly language**, a symbolic way of specifying machine code
 - write instructions using names rather than bit string
 - refer to registers using either numbers or names
 - allow names (labels) to be associated with memory addresses

Examples MIPS Assembler

```
lw  $t1,address    # reg[t1] = memory[address]
sw  $t3,address    # memory[address] = reg[t3]
                   # address must be 4-byte aligned
la  $t1,address    # reg[t1] = address
lui $t2,const      # reg[t2] = const <<< 16
and $t0,$t1,$t2    # reg[t0] = reg[t1] & reg[t2]
add $t0,$t1,$t2    # reg[t0] = reg[t1] + reg[t2]
                   # add signed 2's complement ints
addi $t2,$t3, 5    # reg[t2] = reg[t3] + 5
                   # add immediate, no sub immediate
mult $t3,$t4       # (Hi,Lo) = reg[t3] * reg[t4]
                   # store 64-bit result in
                   # registers Hi,Lo
seq $t7,$t1,$t2    # reg[t7] = (reg[t1] == reg[t2])
j   label           # PC = label
beq $t1,$t2,label  # PC = label if reg[t1]==reg[t2]
nop               # do nothing
```

9

10

MIPS Architecture

MIPS CPU has

- 32 general purpose registers (32-bit)
- 16/32 floating-point registers (for float/double)
- PC ... 32-bit register (always aligned on 4-byte boundary)
- HI,LO ... for storing results of multiplication and division

Registers can be referred to as \$0 .. \$31 or by symbolic names

Some registers have special uses e.g.

- register \$0 always has value 0, cannot be written
- registers \$1, \$26, \$27 reserved for use by system

More details on following slides ...

MIPS Architecture - Integer Registers

Number	Names	Conventional Usage
0	\$zero	Constant 0
1	\$at	Reserved for assembler
2,3	\$v0,\$v1	Expression evaluation and results of a function
4..7	\$a0..\$a3	Arguments 1-4
8..16	\$t0..\$t7	Temporary (not preserved across function calls)
16..23	\$s0..\$s7	Saved temporary (preserved across function calls)
24,25	\$t8,\$t9	Temporary (preserved across function calls)
26,27	\$k0,\$k1	Reserved for OS kernel
28	\$gp	Pointer to global area
29	\$sp	Stack pointer
30	\$fp	Frame pointer
31	\$ra	Return address (used by function call instruction)

- Except for registers 0 and 31, these uses are only conventions.
- Conventions allow compiled code from different sources to be combined (linked).
- Most of these conventions are irrelevant when you are writing small MIPS assembly code programs.
- But use registers 8..25 for holding values.

Reg	Notes
\$f0..\$f2	hold return value of functions which return floating-point results
\$f4..\$f10	temporary registers; not preserved across function calls
\$f12..\$f14	used for first two double-precision function arguments
\$f16..\$f18	temporary registers; used for expression evaluation
\$f20..\$f30	saved registers; value is preserved across function calls

Notes:

- floating point registers can be used as 32 32-bit register or 16 64-bit registers
- for 64-bit use even numbered registers

13

14

Data and Addresses

All operations refer to data, either

- in a register
- in memory
- literally (i.e. constant)

Computation operations refer to registers or constants.

Only load/store instructions refer to memory.

To access registers, you can also use \$name

e.g. \$zero == \$0, \$t0 == \$8, \$fp == \$30, ...

To refer to literals, use C-like constants:

```
1 3 -1 -2 12345 0x1 0xFFFFFFFF
"a string" 'a' 'b' '1' '\n' '\0'
```

Describing MIPS Assembler Operations

- Registers are denoted:

R_d & destination register {(where result goes)} \ R_s & source register #1 {(where data comes from)} \ R_t & source register #2 {(where data comes from)} \

- {Reg[]-{R}{}} = contents of register {R}.
- Data transfer is denoted by {<-}.

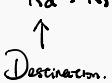
```
add $6, $7, $8 # Reg[6] <- Reg[7] + Reg[8]
```

15

16

<code>li R_d, imm</code>	load immediate
	$\text{Reg}[R_d] \leftarrow \text{imm}$
<code>la R_d, addr</code>	load address
	$\text{Reg}[R_d] \leftarrow \text{addr}$

<code>move R_d, R_s</code>	move data reg-to-reg
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s]$

Remember $R_d = R_s$


- The `{li}` (*load immediate*) instruction is used to set a register to a constant value, e.g

```
li $8, 42      # $8 = 42
li $24, 0x2a   # $24 = 42
li $15, '*'    # $15 = 42
```

- The `{move}` instruction is used to set a register to the same value as another register, e.g

```
move $8, $9    # assign to $8 value in $9
```

- Note destination is first register.

17

18

Setting A Register to An Address

Pseudo Instructions

- Note the `{la}` (*load address*) instruction is used to set a register to a labelled memory address.

```
la $8, start
```

- The memory address will be fixed before the program is run, so this differs only syntactically from the `{li}` instruction.

- For example, if `vec` is the label for memory address `0x10000100` then these two instructions are equivalent:

```
la $7, vec
li $7, 0x10000100
```

- In both cases the constant is encoded as part of the instruction.

- Neither `{la}` or `{li}` access memory - there are very different to the `{lw}` instruction.

- Both `la` and `li` are pseudo instructions provided by the assembler for user convenience.
- The assembler translates these pseudo-instructions into instructions actually implemented by the processor.
- For example, `li $7, 15` might be translated to `addi $7, $0, 15`
- If the constant is large the assembler will need to translate a `li/la` instruction to two actual instructions.

19

20

Arithmetic Instructions

<code>add R_d, R_s, R_t</code>	<i>addition</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] + \text{Reg}[R_t]$
<code>add R_d, R_s, imm</code>	<i>addition</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] + \text{imm}$
<code>sub R_d, R_s, R_t</code>	<i>subtraction</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] - \text{Reg}[R_t]$
<code>mul R_d, R_s, R_t</code>	<i>multiplication</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] * \text{Reg}[R_t]$
<code>div R_d, R_s, R_t</code>	<i>division</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] / \text{Reg}[R_t]$
<code>rem R_d, R_s, R_t</code>	<i>remainder</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] \% \text{Reg}[R_t]$
<code>neg R_d, R_s</code>	<i>negate</i>
	$\text{Reg}[R_d] \leftarrow -\text{Reg}[R_s]$

All arithmetic is signed (2's-complement).

Logic Instructions

<code>and R_d, R_s, R_t</code>	<i>logical and</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] \& \text{Reg}[R_t]$
<code>and R_d, R_s, imm</code>	<i>logical and</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] \& \text{imm}$
<code>or R_d, R_s, R_t</code>	<i>logical or</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] \text{Reg}[R_t]$
<code>not R_d, R_s</code>	<i>logical not</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s]$
<code>xor R_d, R_s, R_t</code>	<i>logical xor</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] ^ \text{Reg}[R_t]$

All of these instructions can use *imm* instead of *R_t*.

Bit Manipulation Instructions

<code>sll R_d, R_s, R_t</code>	<i>shift left logical</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] << \text{Reg}[R_t]$
<code>sll R_d, R_s, imm</code>	<i>shift left logical</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] << \text{imm}$
<code>srl R_d, R_s, imm</code>	<i>shift right logical</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] >> \text{imm}$
<code>sra R_d, R_s, imm</code>	<i>shift right arithmetic</i>
	$\text{Reg}[R_d] \leftarrow \text{Reg}[R_s] >> \text{imm}$
<code>rol R_d, R_s, imm</code>	<i>rotate left</i>
	$\text{Reg}[R_d] \leftarrow \text{rot}(\text{Reg}[R_s], \text{imm}, \text{left})$
<code>ror R_d, R_s, imm</code>	<i>rotate right</i>
	$\text{Reg}[R_d] \leftarrow \text{rot}(\text{Reg}[R_s], \text{imm}, \text{right})$

All of these instructions can use *R_t* instead of *imm*.

Jump Instructions

Jumps control flow of program execution.

Branch Instructions

Branches combine testing and jumping.

j *label* jump to location
PC <- & label
jal *label* jump and link

jr *R_s* jump via register
PC <- Reg[*R_s*]
jalr *R_s* jump and link via reg

beq <i>R_s, R_t, label</i>	branch on equal if (Reg[<i>R_s</i>] == Reg[<i>R_t</i>]) PC <- label
bne <i>R_s, R_t, label</i>	branch on not equal if (Reg[<i>R_s</i>] != Reg[<i>R_t</i>]) PC <- label
blt <i>R_s, R_t, label</i>	branch on less than if (Reg[<i>R_s</i>] < Reg[<i>R_t</i>]) PC <- label
ble <i>R_s, R_t, label</i>	branch on less or equal if (Reg[<i>R_s</i>] <= Reg[<i>R_t</i>]) PC <- label
bgt <i>R_s, R_t, label</i>	branch on greater than if (Reg[<i>R_s</i>] > Reg[<i>R_t</i>]) PC <- label
bge <i>R_s, R_t, label</i>	branch on greater or equal if (Reg[<i>R_s</i>] >= Reg[<i>R_t</i>]) PC <- label

25

26

MIPS Instruction Set

Implementation of pseudo-instructions:

What you write	Machine code produced
li \$t5, const	ori \$t5, \$0, const
la \$t3, label	lui \$at, label[31..16] <i>Don't get me</i> ori \$t3, \$at, label[15..0]
bge \$t1, \$t2, label	slt \$at, \$t1, \$t2 <i>if (\$t1 < \$t2) { } </i> beq \$at, \$0, label
blt \$t1, \$t2, label	slt \$at, \$t1, \$t2 bne \$at, \$0, label

Note: use of \$at register for intermediate results

MIPS vs SPIM

MIPS is a machine architecture, including instruction set

SPIM is an *emulator* for the MIPS instruction set

- reads text files containing instruction + directives
- converts to machine code and loads into "memory"
- provides debugging capabilities
 - single-step, breakpoints, view registers/memory, ...
- provides mechanism to interact with operating system (syscall)

Also provides extra instructions, mapped to MIPS core set

- provide convenient/mnemonic ways to do common operations
- e.g. move \$s0,\$v0 rather than addu \$s0,\$0,\$v0

27

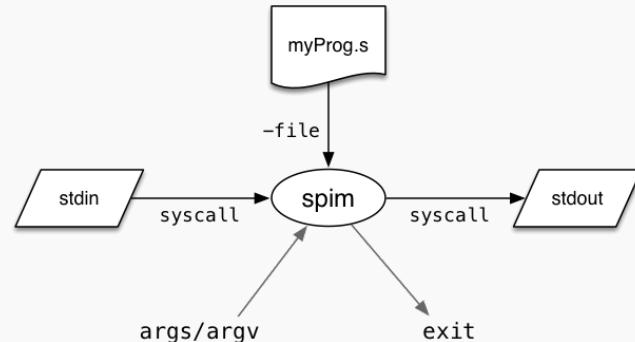
28

Three ways to execute MIPS code with SPIM

✓ spim ... command line tool

- load programs using -file option
 - interact using stdin/stdout via login terminal
- qtspim ... GUI environment
- load programs via a load button
 - interact via a pop-up stdin/stdout terminal
- xspim ... GUI environment
- similar to qtspim, but not as pretty
 - requires X-windows server

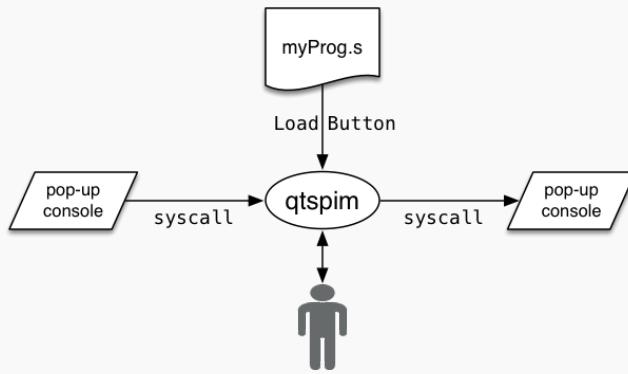
Command-line tool:



29

30

GUI tool:



```

$ 1521 spim
Loaded: /home/cs1521/share/spim/exceptions.s
(spim) load "myprogram.s"
(spim) step 6
[0x00400000] 0x8fa40000 lw $4, 0($29)
[0x00400004] 0x27a50004 addiu $5, $29, 4
[0x00400008] 0x24a60004 addiu $6, $5, 4
[0x0040000c] 0x00041080 sll $2, $4, 2
[0x00400010] 0x00c23021 addu $6, $6, $2
[0x00400014] 0x0c100009 jal 0x00400024 [main]
(spim) print_all_regs hex
...
General Registers
R0 (r0) = 00000000 R8 (t0) = 00000000 R16 (s0) = 00000000 ...
R1 (at) = 10010000 R9 (t1) = 00000000 R17 (s1) = 00000000 ...
  
```

31

32

The SPIM interpreter provides I/O and memory allocation via the syscall instruction.

Service	n	Arguments	Result
printf("%d")	1	int in \$a0	-
printf("%f")	2	float in \$f12	-
printf("%lf")	3	double in \$f12	-
printf("%s")	4	\$a0 = string	-
scanf("%d")	5	-	int in \$v0
scanf("%f")	6	-	float in \$f0
scanf("%lf")	7	-	double in \$f0
fgets	8	buffer address in \$a0	
		length in \$a1	-
sbrrk	9	nbytes in \$a0	address in \$v0
printf("%c")	11	char in \$a0	-
scanf("%c")	12	-	char in \$v0
exit(status)	17	status in \$a0	-

33

Region	Address	Notes
text	0x00400000	instructions only; read-only; cannot expand
data	0x10000000	data objects; read/write; can be expanded
stack	0x7fffffff	grows down from that address; read/write
k_text	'0x80000000'	kernel code; read-only only accessible in kernel mode
k_data	'0x90000000'	kernel data' only accessible in kernel mode

34

MIPS Assembly Language

MIPS assembly language programs contain

- comments ... introduced by #
- labels ... appended with :
- directives ... symbol beginning with .
- assembly language instructions

Programmers need to specify

- data objects that live in the data region
- functions (instruction sequences) that live in the code/text region

Each instruction or directive appears on its own line

Example MIPS assembler program

```
# hello.s ... print "Hello, MIPS"

main:
    la $a0, msg      # load the argument string
    li $v0, 4        # load the system call (print)
    syscall          # print the string
    jr $ra           # return to caller (_start)

.data            # the data segment
msg: .asciiz "Hello, MIPS\n"
```

```
# Prog.s ... comment giving description of function
# Author ...

main:      # indicates start of code
           # (i.e. first user instruction to execute)
           # ...

.data       # variable declarations follow this line
           # ...

# End of program; leave a blank line to make SPIM happy
```

Directives {(instructions to assembler, not MIPS instructions)}

```
.text      # following instructions placed in text
.data      # following objects placed in data

.globl     # make symbol available globally

a: .space 18 # uchar a[18]; or uint a[4];
           # align next object on 2-byte addr  $2^{\circ} \cdot 4$ 
           the power of:

i: .word 2   # unsigned int i = 2;
v: .word 1,3,5 # unsigned int v[3] = {1,3,5};
h: .half 2,4,6 # unsigned short h[3] = {2,4,6};
b: .byte 1,2,3 # unsigned char b[3] = {1,2,3};
f: .float 3.14 # float f = 3.14;

s: .asciiz "abc" # char s[4] {'a','b','c','\0'};
t: .ascii "abc" # char s[3] {'a','b','c'};
```

37

38

Encoding MIPS Instructions as 32 bit Numbers

Assembler	Encoding
\$ 7 = \$ 8 + \$ 0	000000 sssss ttttt dddddd00000100000
add \$ d, \$ s, \$ t	000000 00111 01000 0000000000100000
add \$ 7, \$ 8, \$ 0	0x01e80020 == 31981600
\$ 5 = \$ 1 - \$ 3	000000 sssss ttttt dddddd00000100010
sub \$ d, \$ s, \$ t	000000 00001 00011 0010100000100010
sub \$ 5, \$ 1, \$ 3	0x00232822 == 2304034
\$ 2 = \$ 2 + 1	001000 sssss dddddd CCCCCCCCCCCCCCCC
addi \$ d, \$ s, C	001000 00010 00010 0000000000000001
addi \$ 2, \$ 2, 1	0x20420001 == 541196289

All instructions are variants of 3 patterns of bits.

Which simplifies silicon and human decoding (ass1!).

39

The {goto} statement allows transfer of control to any labelled point with a function. For example, this code:

```
for (int i = 1; i <= 10; i++) {
    printf("%d\n", i);
}
```

can be written as:

```
int i = 1;
loop:
    if (i > 10) goto end;
    i++;
    printf("%d", i);
    printf("\n");
    goto loop;
end:
```

- goto statements can result in very difficult to read programs.
- goto statements can also result in slower programs.
- In general, use of goto is considered poor programming style.
- Do not use goto without very good reason.
- kernel & embedded programmers sometimes use goto.

MIPS Programming

Writing correct assembler directly is hard.

Recommended strategy:

- develop the solution in C
- map to “simplified” C
- translate each simplified C statement to MIPS instructions

Simplified C

- does *not* have while, compound if, complex expressions
- *does* have simple if, goto, one-operator expressions

Simplified C makes extensive use of

- *labels* ... symbolic name for C statement
- *goto* ... transfer control to labelled statement

Example:

Mapping C into MIPS

Things to do:

- allocate variables to registers/memory
- place literals in data segment
- transform C program to:
 - break expression evaluation into steps
 - replace control structures by goto

add: C to simplified C

```
int main(void) {  
    int x = 17;  
    int y = 25;  
    printf("%d\n", x + y);  
}
```

```
int main(void) {  
    int x, y, z;  
    x = 17;  
    y = 25;  
    z = x + y;  
    printf("%d", z);  
    printf("\n");  
}
```

add: simplified C to MIPS

```
int main(void) {  
    int x, y, z;  
    x = 17;  
    y = 25;  
    z = x + y;  
    printf("%d", z);  
    printf("\n");  
}  
  
main:  
    li    $t0, 17  
    li    $t1, 25  
    add  $t2, $t1, $t0  
    move $a0, $t2  
    li    $v0, 1  
    syscall  
  
    li    $a0, '\n'  
    li    $v0, 11  
    syscall  
  
    jr   $ra
```

5

6

while: C to simplified C

```
i = 0;  
n = 0;  
while (i < 5) {  
    n = n + i;  
    i++;  
}  
  
i = 0;  
n = 0;  
loop:  
    if (i >= 5) goto end;  
    n = n + i;  
    i++;  
    goto loop;  
end:
```

while: simplified C to MIPS

```
i = 0;  
n = 0;  
loop:  
    if (i >= 5) goto end;  
    n = n + i;  
    i++;  
    goto loop;  
end:  
  
li $t0, 0 # i in $t0  
li $t1, 0 # n in $t1  
loop:  
    bge $t0, 5, end  
    add $t1, $t1, $t0  
    add $t0, $t0, 1  
    goto loop  
end:
```

7

8

if: C to simplified C

```
if (i < 0) {  
    n = n - i;  
} else {  
    n = n + i;  
}  
  
if (i >= 0) goto else1;  
n = n - i;  
goto end1;  
else1:  
    n = n + i;  
end1:  
{Note: you can't use else as a {label} in C}
```

if: simplified C to MIPS

```
# assume i in $t0  
# assume n in $t1  
bge $t0, 0, else1  
sub $t1, $t1, $t0  
goto end1  
else1:  
    add $t1, $t1, $t0  
end1:
```

9

10

if/and: C to simplified C

```
if (i < 0 && n >= 42) {  
    n = n - i;  
} else {  
    n = n + i;  
}  
  
if (i >= 0) goto else1;  
if (n < 42) goto else1;  
n = n - i;  
goto end1;  
else1:  
    n = n + i;  
end1:
```

if/and: simplified C to MIPS

```
# assume i in $t0  
# assume n in $t1  
bge $t0, 0, else1  
blt $t1, 42, else1  
sub $t1, $t1, $t0  
goto end1  
else1:  
    add $t1, $t1, $t0  
end1:
```

11

12

```

if (i < 0 || n >= 42) {
    n = n - i;
} else {
    n = n + i;
}
if (i < 0) goto then1;
if (n >= 42) goto then1;
goto else1;
then1:
    n = n - i;
    goto end1;
else1:
    n = n + i;
end1:

```

```

int main(void) {
    for (int i = 0; i <= 10; i++) {
        printf("%d\n", i);
    }
}

```

13

14

Example Printing First 10 Integers

Convert to goto and simple C statements and decide where variables will be stored.

```

int main(void) {
    int i; // in register $t0
    i = 0;
loop:
    if (i >= 10)
        goto end;
    i++;
    printf("%d", i);
    printf("\n");
    goto loop;
end:
    return 0;
}

```

Example Printing First 10 Integers

```

main:                      # int main(void) {
                            # int i; // in register $t0
    li    $t0, 0            # i = 0;
loop:                      # loop:
    bge  $t0, 10 end     # if (i >= 10) goto end;
    add  $t0, $t0 1        # i++;
    move $a0, $t0          # printf("%d" i);
    li    $v0, 1
    syscall
    li    $a0, '\n'        # printf("%c", '\n');
    li    $v0, 11
    syscall
    b    loop             # goto loop;
end:
    jr   $ra              # return

```

15

16

Directives all(instructions to assembler, not MIPS instructions)}

```
.text      # following instructions placed in text
.data      # following objects placed in data
.globl    # make symbol available globally

a: .space 18 # uchar a[18]; or uint a[4];
   .align 2  # align next object on 2-byte addr

i: .word 2   # unsigned int i = 2;
v: .word 1,3,5 # unsigned int v[3] = {1,3,5};
h: .half 2,4,6 # unsigned short h[3] = {2,4,6};
b: .byte 1,2,3 # unsigned char b[3] = {1,2,3};
f: .float 3.14 # float f = 3.14;

s: .asciiz "abc" # char s[4] {'a','b','c','\0'};
t: .ascii "abc" # char s[3] {'a','b','c'};
```

1

Region	Address	Notes
text	0x00400000	instructions only; read-only; cannot expand
data	0x10000000	data objects; read/write; can be expanded
stack	0x7fffffff	grows down from that address; read/write
k_text	'0x80000000'	kernel code; read-only only accessible in kernel mode
k_data	'0x90000000'	kernel data' only accessible in kernel mode

2

Setting A Register to An Address

Accessing Memory

- Note the `la` (load address) instruction is used to set a register to a labelled memory address.

`la $8, start`

- The memory address will be fixed before the program is run, so this differs only syntactically from the `{li}` instruction.

- For example, if `vec` is the label for memory address `0x10000100` then these two instructions are equivalent:

```
la $7, vec
li $7, 0x10000100
```

- In both cases the constant is encoded as part of the instruction(s).
- Neither `la` or `li` access memory - there are very different to the `lw` instruction.

- only load/store instructions move data between memory and CPU.
- 1, 2 and 4-bytes (8, 16 and 32 bit) quantities can be moved.
 - two bytes ("halfword"), four bytes ("word")
- Two operands: the register which supplies/receives the value and the memory address.
- for 1 and 2-byte operations the low (least significant) bits of the register are used.
- `LB` & `LH` assume bytes/halfword contain a 8-bit/16-bit signed integer
- unsigned equivalents `LHU` and `LBU` assume integer is unsigned

Accessing Memory

assembly	meaning	bit pattern
lb $r_t, I(r_b)$	$r_t = \text{mem}[r_b+I]$	100000bbbbbtttt000000000000000000000000
lh $r_t, I(r_b)$	$r_t = \text{mem}[r_b+I] \text{mem}[r_b+I+1] \ll 8$	100001bbbbbtttt000000000000000000000000
lw $r_t, I(r_b)$	$r_t = \text{mem}[r_b+I] \text{mem}[r_b+I+1] \ll 8 \text{mem}[r_b+I+2] \ll 16 \text{mem}[r_b+I+3] \ll 24$	100011bbbbbtttt000000000000000000000000
sb $r_t, I(r_b)$	$\text{mem}[r_b+I] = r_t \& 0xff$	101000bbbbbtttt000000000000000000000000
sh $r_t, I(r_b)$	$\text{mem}[r_b+I] = r_t \& 0xff$ $\text{mem}[r_b+I+1] = r_t \gg 8 \& 0xff$	101001bbbbbtttt000000000000000000000000
sw $r_t, I(r_b)$	$\text{mem}[r_b+I] = r_t \& 0xff$ $\text{mem}[r_b+I+1] = r_t \gg 8 \& 0xff$ $\text{mem}[r_b+I+2] = r_t \gg 16 \& 0xff$ $\text{mem}[r_b+I+3] = r_t \gg 24 \& 0xff$	101011bbbbbtttt000000000000000000000000

Memory Addressing Modes

Ways of specifying memory addresses:

Format	Address referred to
(reg)	contents of register e.g. (\$8) (\$fp) (\$5)
imm	immediate (= constant) e.g. 0x80000000 123456789
imm(reg)	immediate + contents of register e.g. 0(\$fp) 0x80000000(\$9)
sym	address of symbol (= name) e.g. main endloop exit
sym(reg)	address of symbol + reg contents e.g. main(\$fp) array(\$9)
sym +/- imm	address of symbol +/- immediate e.g. main+8 endloop-4
sym +/- imm (reg)	sym address +/- (imm + reg contents) e.g. main+8(\$8) array+0(\$18)

5

6

Addressing Modes

Examples of load/store and addressing:

```
main:
    la $t0, vec      # reg[t0] = &vec[0]
    li $t1, 5        # reg[t1] = 5
    sw $t1, ($t0)    # vec[0] = reg[t1]
    li $t1, 13       # reg[t1] = 13
    sw $t1, 4($t0)   # vec[1] = reg[t1]
    li $t1, -7       # reg[t1] = -7
    sw $t1, 8($t0)   # vec[2] = reg[t1]
    li $t2, 12       # reg[t2] = 12
    li $t1, 42       # reg[t1] = 42
    sw $t1, vec($t2) # vec[3] = reg[t1]
    jr $ra            # return
.data
vec: .space 16      # 16 bytes of storage
```

Data Structures and MIPS

C data structures and their MIPS representations:

- char ... as byte in memory, or register
- int ... as 4 bytes in memory, or register
- double ... as 8 bytes in memory, or \$f? register
- arrays ... sequence of bytes in memory, elements accessed by index (calculated on MIPS)
- structs ... sequence of bytes in memory, accessed by fields (constant offsets on MIPS)

A char, int or double

- can be stored in register if local variable and no pointer to it
- otherwise stored on stack if local variable
- stored in data segment if global variable

7

8

- global/static variables need appropriate number of bytes allocated in data segment using {.space}:

```
double val;           val: .space 8
char str[20];         str: .space 20
int vec[20];          vec: .space 80
```

initialized to 0 by default, other directives allow initialization to other values:

```
int val = 5;           val: .double 5
int arr[4] = {9,8,7,6}; arr: .word 9, 8, 7, 6
char msg[7] = "Hello\n"; msg: .ascii "Hello\n"
```

```
C
int main(void) {
    int x, y, z;
    x = 17;
    y = 25;
    z = x + y;
```

```
MIPS
main:
# x in $t0
# y in $t1
# z in $t2
li   $t0, 17
li   $t1, 25
add $t2, $t1, $t0
// ...
```

9

10

add: variables in memory

store value in array element - v1

```
C
int x, y, z;
int main(void) {
    x = 17;
    y = 25;
    z = x + y;
}

.data
x: .space 4
y: .space 4
z: .space 4
```

```
MIPS
main:
li   $t0, 17
sw   $t0, x
li   $t0, 25
sw   $t0, y
lw   $t0, x
lw   $t1, y
add $t2, $t1, $t0
sw   $t2, z
```

```
C
int x[10];
int main(void) {
    // sizeof x[0] == 4
    x[3] = 17;
}
```

```
MIPS
main:
li   $t0, 3
# each array element
# is 4 bytes
mul $t0, $t0, 4
la   $t1, x
add $t2, $t1, $t0
li   $t3, 17
sw   $t3, ($t2)

.data
x: .space 40
```

11

12

```
C
#include <stdint.h>

int16_t x[30];

int main(void) {
    // sizeof x[0] == 2
    x[13] = 23;
}
```

```
MIPS
main:
    li    $t0, 13
    # each array element
    # is 2 bytes
    mul  $t0, $t0, 2
    la   $t1, x
    add  $t2, $t1, $t0
    li   $t3, 23
    sw   $t3, ($t2)
.data
x:   .space 60
```

Can be named initialised as noted above:

```
vec: .space 40
# could be either int vec[10] or char vec[40]

nums: .word 1, 3, 5, 7, 9
# int nums[6] = {1,3,5,7,9}
```

Can access elements via index or cursor (pointer)

- either approach needs to account for size of elements

Arrays passed to functions via pointer to first element

- must also pass array size, since not available elsewhere

{See sumOf() exercise for an example of passing an array to a function }

13

14

Printing 1-d Arrays in MIPS - v1

```
C
int vec[5]={0,1,2,3,4};
// ...
int i = 0
while (i < 5) {
    printf("%d", vec[i]);
    i++;
}
// ....
■ i in $s0
```

```
MIPS
li    $s0, 0
loop:
    bge  $s0, 5, end
    la   $t0, vec
    mul  $t1, $s0, 4
    add  $t2, $t1, $t0
    lw   $a0, ($t2)
    li   $v0, 1
    syscall
    addi $s0, $s0, 1
    b    loop
end:
.data
vec: .word 0,1,2,3,4
```

Printing 1-d Array in MIPS - v2

```
C
int vec[5]={0,1,2,3,4};
// ...
int *p = &vec[0];
int *end = &vec[4];
while (p <= end) {
    int y = *p;
    printf("%d", y);
    p++;
}
// ....
■ p in $s0
■ end in $s1
```

```
MIPS
li    $s0, vec
la   $t0, vec
add  $s1, $t0, 16
loop:
    bgt $s0, $s1, end
    lw   $a0, ($s0)
    li   $v0, 1
    syscall
    addi $s0, $s0, 4
    b    loop
end:
.data
vec: .word 0,1,2,3,4
```

15

16

1-d Arrays in MIPS

Scanning across an array of N elements using cursor

```
# int vec[10] = {...};  
# int *cur, *end = &vec[10];  
# for (cur = vec; cur < end; cur++)  
#   printf("%d\n", *cur);}  
  
la  $s0, vec           # cur = &vec[0]  
la  $s1, vec+40        # end = &vec[10]  
loop:  
  bge $s0, $s1, end_loop # if (cur >= end) break  
  lw   $a0, ($s0)          # a0 = *cur  
  jal  print             # print a0  
  addi $s0, $s0, 4         # cur++  
  j    loop  
end_loop:
```

{Assumes the existence of a print() function to do printf("\%d n",x)}

2-d Arrays in MIPS

Representations of int matrix[4][4] ...

```
matrix: .space 64
```

Now consider summing all elements

```
int i, j, sum = 0;  
for (i = 0; i < 4; i++) {  
  for (j = 0; j < 4; j++) {  
    sum += matrix[i][j];  
  }  
}
```

2-d Arrays in MIPS

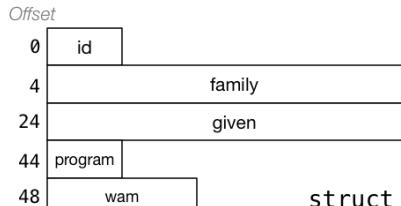
Computing sum of all elements in int matrix[6][5] in C

```
int row, col, sum = 0;  
  
// row-by-row  
for (row = 0; row < 6; row++) {  
  // col-by-col within row  
  for (col = 0; col < 5; row++) {  
    sum += matrix[row][col];  
  }  
}
```

2-d Arrays in MIPS

Computing sum of all elements int matrix[6][5]

```
li  $s0, 0           # sum = 0  
li  $s1, 6           # s1 = #rows  
li  $s2, 0           # row = 0  
li  $s3, 5           # s3 = #cols  
li  $s4, 0           # col = 0 // redundant  
li  $s5, 4           # intsize = sizeof(int)  
mul $s6, $s3, $s5   # rowsize = #cols*intsize  
loop1:  
  bge $s2, $s1, end1 # if (row >= 6) break  
  li   $s4, 0          # col = 0  
loop2:  
  bge $s4, $s3, end2 # if (col >= 5) break  
  mul $t0, $s2, $s6    # t0 = row*rowsize  
  mul $t1, $s4, $s5    # t1 = col*intsize  
  add $t0, $t0, $t1    # offset = t0+t1  
  lw   $t0, matrix($t0) # t0 = *(matrix+offset)
```



C struct definitions effectively define a new type.

```
// new type called "struct student"
struct student {...};

// new type called student_t
typedef struct student student_t;
```

Instances of structures can be created by allocating space:

```
# sizeof(Student) == 56
stu1: # student_t stu1;
       .space 56
stu2: # student_t stu2;
       .space 56
stu:   # student_t *stu;
       .space 4
```

Structs in MIPS

Accessing structure components is by offset, not name

```

stu1: .space 56      # student_t stu1;
stu2: .space 56      # student_t stu2;
# stu is $s1          # student_t *stu;

li $t0 5012345
sw $t0, stu1+0      # stu1.id = 5012345;
li $t0, 3778
sw $t0, stu1+44      # stu1.program = 3778;

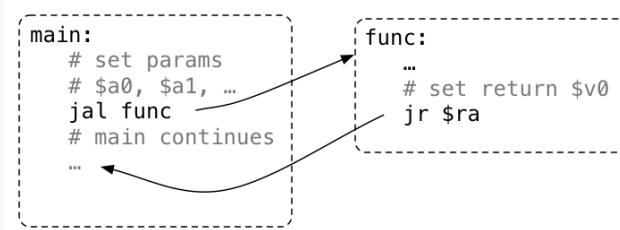
la $s1, stu2        # stu = &stu2;
li $t0, 3707
sw $t0, 44($s1)      # stu->program = 3707;
li $t0, 5034567
sw $t0, 0($s1)      # stu->id = 5034567;
```

When we call a function:

- the arguments are evaluated and set up for function
- control is transferred to the code for the function
- local variables are created
- the function code is executed in this environment
- the return value is set up
- control transfers back to where the function was called from
- the caller receives the return value

Simple view of function calls:

- load argument values into \$a0, \$a1, \$a2, \$a3.
- jal function set \$ra to PC+4 and jumps to function
- function puts return value in \$v0
- returns to caller using jr \$ra



1

2

Function with No Parameters or Return Value

Function with a Return Value but No Parameters

- jal hello sets \$ra to address of following instruction and transfers execution to hello
- jr \$ra transfers execution to the address in \$ra

```

int main(void) {
    main:
    hello();
    ...
    return 0;
}
}

void hello(void) {
    hello:
    ...
    la $a0, string
    li $v0, 4
    syscall
    jr $ra
    .data
    string:
    .asciiz "hi\n"
}

```

- by convention function return value is passed back in \$v0

```

int main(void) {
    int a = answer();
    printf("%d\n", a);
    return 0;
}

int answer(void) {
    return 42;
}

```

3

4

Function with a Return Value and Parameters

Function calling another function - DO NOT DO THIS

- by convention first 4 function parameters passed in \$a0 \$a1 \$a2 \$a3
- if there are more parameters they are passed on the stack

```
int main(void) {                                main:  
    int a = product(6, 7);                      ...  
    printf("%d\n", a);                          li  $a0, 6  
    return 0;                                    li  $a1, 7  
}                                              jal  product  
                                                move $a0, $v0  
int product(int x, int y) {                  li  $v0, 1  
    return x * y;                            syscall  
}                                              ...  
                                                product:  
                                                mul  $v0, $a0, $a1  
                                                jr   $ra
```

- a function that calls another function must save \$ra
- the jr \$ra in main below will fail because jal hello changed \$ra

```
int main(void) {                                main:  
    hello();                                 jal  hello  
    return 0;                               li  $v0, 0  
}                                              jr  $ra # THIS WILL FAIL  
hello:  
void hello(void) {                           la  $a0, string  
    printf("hi\n");                         li  $v0, 4  
}                                              syscall  
                                                jr   $ra  
                                                .data  
string: .asciiz "hi\n"
```

5

6

Simple Function Call Example - C

```
void f(void);  
int main(void) {  
    printf("calling function f\n");  
    f();  
    printf("back from function f\n");  
    return 0;  
}  
void f(void) {  
    printf("in function f\n");  
}
```

source code for call_return.c

Simple Function Call Example - broken MIPS

```
la  $a0, string0  # printf("calling function f\n");  
li  $v0, 4  
syscall  
jal f            # set $ra to following address  
la  $a0, string1  # printf("back from function f\n");  
li  $v0, 4  
syscall  
li  $v0, 0          # fails because $ra changes since main call  
jr  $ra            # return from function main  
f:  
la  $a0, string2  # printf("in function f\n");  
li  $v0, 4  
syscall  
jr  $ra            # return from function f  
.data
```

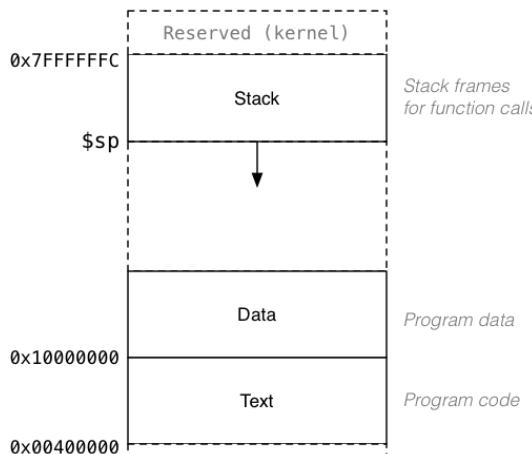
source code for call_return.broken.s

7

8

Stack - Where it is in Memory

Data associated with a function call placed on the stack:



Stack - Using stack to Save/Restore registers

```
f:
    sub $sp, $sp, 12      # allocate 12 bytes
    sw $ra, 8($sp)        # save $ra on $stack
    sw $s1, 4($sp)        # save $s1 on $stack
    sw $s0, 0($sp)        # save $s0 on $stack

    ...

    lw $s0, 0($sp)        # restore $s0 from $stack
    lw $s1, 4($sp)        # restore $s1 from $stack
    lw $ra, 8($sp)        # restore $ra from $stack
    add $sp, $sp, 12       # move stack pointer back
    jr $ra                 # return
```

Stack - Allocating Space

- \$sp (stack pointer) initialized by operating system
- always 4-byte aligned (divisible by 4)
- points at currently used (4-byte) word
- grows downward
- a function can do this to allocate 40 bytes:

```
sub $sp, $sp, 40      # move stack pointer down
```

- a function **must** leave \$sp at original value
- so if you allocated 40 bytes, before return (jr \$ra)

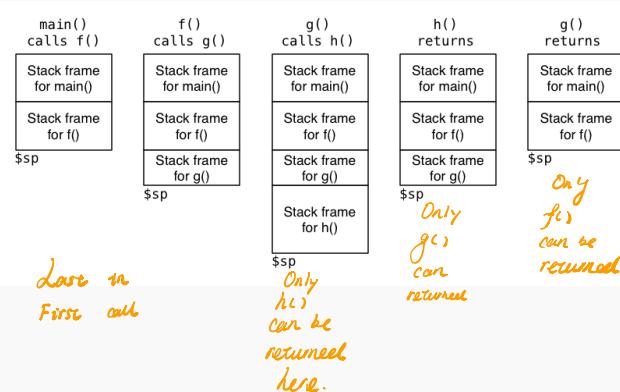
```
add $sp, $sp, 40      # move stack pointer back
```

9

10

Stack - Growing & Shrinking

How stack changes as functions are called and return:



11

12

Function calling another function - How to Do It

- a function that calls another function must save \$ra

```
main:  
    sub $sp, $sp, 4      # move stack pointer down  
                        # to allocate 4 bytes  
    sw $ra, 0($sp)      # save $ra on $stack  
  
    jal hello           # call hello  
  
    lw $ra, 0($sp)      # recover $ra from $stack  
    add $sp, $sp, 4      # move stack pointer back up  
                        # to what it was when main called  
    li $v0, 0            # return 0  
    jr $ra               #
```

Simple Function Call Example - correct MIPS

```
la $a0, string0  # printf("calling function f\n");  
li $v0, 4  
syscall  
jal f             # set $ra to following address  
la $a0, string1  # printf("back from function f\n");  
li $v0, 4  
syscall  
lw $ra, 0($sp)   # recover $ra from $stack  
addi $sp, $sp, 4  # move stack pointer back to what it was  
li $v0, 0          # return 0 from function main  
jr $ra             #  
f:  
    la $a0, string2  # printf("in function f\n");  
    li $v0, 4  
    syscall  
    jr $ra             # return from function f
```

13

source code for call_return.s

MIPS Register usage conventions

- a0..a3 contain first 4 arguments
- \$v0 contains return value
- \$ra contains return address
- if function changes \$sp, \$fp, \$s0..\$s8 it restores their value
- callers assume \$sp, \$fp, \$s0..\$s8 unchanged by call (jal)
- a function may destroy the value of other registers e.g. \$t0..\$t9
- callers must assume value in e.g. \$t0..\$t9 changed by call (jal)

Can I use counter?

MIPS Register usage conventions - not covered in COMP1521

- floating point registers used to pass/return float/doubles
- similar conventions for saving floating point registers
- stack used to pass arguments after first 4
- stack used to pass arguments which do not fit in register
- stack used to return values which do not fit in register
- for example C argument or return value can be a struct, which is any number of bytes

15

16

Example - Returning a Value - C

```
int answer(void);
int main(void) {
    int a = answer();
    printf("%d\n", a);
    return 0;
}

int answer(void) {
    return 42;
}
```

source code for return_answer.c

Example - Returning a Value - MIPS

```
main:
    addi $sp, $sp, -4 # move stack pointer down to make room
    sw   $ra, 0($sp)  # save $ra on $stack
    jal  answer        # call answer, return value will be in $v0
    move $a0, $v0       # printf("%d", a);
    li   $v0, 1
    syscall
    li   $a0, '\n'     # printf("%c", '\n');
    li   $v0, 11
    syscall
    lw   $ra, 0($sp)  # recover $ra from $stack
    addi $sp, $sp, 4   # move stack pointer back up to what it was
    jr   $ra            #
answer: # code for function answer
    li   $v0, 42        #
    jr   $ra            # return from answer
```

17

source code for return_answer.s

18

Example - Argument & Return - C

```
void two(int i);
int main(void) {
    two(1);
}
void two(int i) {
    if (i < 1000000) {
        two(2 * i);
    }
    printf("%d\n", i);
}
```

source code for two_powerful.c

Example - Argument & Return - MIPS (main)

```
main:
    addi $sp, $sp, -4   # move stack pointer down to make room
    sw   $ra, 0($sp)   # save $ra on $stack
    li   $a0, 1         # two(1);
    jal  two           #
    lw   $ra, 0($sp)   # recover $ra from $stack
    addi $sp, $sp, 4   # move stack pointer back up to what it was
    jr   $ra            # return from function main
```

source code for two_powerful.s

19

20

Example - Argument & Return - MIPS (two)

```
two:  
    addi $sp, $sp, -8    # move stack pointer down to make room  
    sw   $ra, 4($sp)      # save $ra on $stack  
    sw   $a0, 0($sp)      # save $a0 on $stack  
    bge $a0, 1000000, print  
    mul $a0, $a0, 2      # restore $a0 from $stack  
    jal two  
  
print:  
    lw   $a0, 0($sp)      # restore $a0 from $stack  
    li   $v0, 1             # printf("%d");  
    syscall  
    li   $a0, '\n'         # printf("%c", '\n');  
    li   $v0, 11  
    syscall  
    lw   $ra, 4($sp)      # restore $ra from $stack  
    addi $sp, $sp, 8        # move stack pointer back up to what it was  
    jr   $ra               # return from two
```

21

Example - more complex Calls - MIPS (main)

```
main:  
    addi $sp, $sp, -4    # move stack pointer down to make room  
    sw   $ra, 0($sp)      # save $ra on $stack  
    li   $a0, 10           # sum_product(10, 12);  
    li   $a1, 12  
    jal  sum_product  
    move $a0, $v0          # printf("%d", z);  
    li   $v0, 1  
    syscall  
    li   $a0, '\n'         # printf("%c", '\n');  
    li   $v0, 11  
    syscall  
    lw   $ra, 0($sp)      # recover $ra from $stack  
    addi $sp, $sp, 4        # move stack pointer back up to what it was  
    li   $v0, 0             # return 0 from function main  
    jr   $ra               # return from function main
```

source code for more_calls.s

Example - More complex Calls - C

```
int main(void) {  
    int z = sum_product(10, 12);  
    printf("%d\n", z);  
    return 0;  
}  
int sum_product(int a, int b) {  
    int p = product(6, 7);  
    return p + a + b;  
}  
int product(int x, int y) {  
    return x * y;  
}
```

source code for more_calls.c

21

22

Example - more complex Calls - MIPS (sum_product)

```
sum_product:  
    addi $sp, $sp, -12   # move stack pointer down to make room  
    sw   $ra, 8($sp)      # save $ra on $stack  
    sw   $a1, 4($sp)      # save $a1 on $stack  
    sw   $a0, 0($sp)      # save $a0 on $stack  
    li   $a0, 6             # product(6, 7);  
    li   $a1, 7  
    jal  product  
    lw   $a1, 4($sp)      # restore $a1 from $stack  
    lw   $a0, 0($sp)      # restore $a0 from $stack  
    add $v0, $v0, $a0       # add a and b to value returned in $v0  
    add $v0, $v0, $a1       # and put result in $v0 to be returned  
    lw   $ra, 8($sp)      # restore $ra from $stack  
    addi $sp, $sp, 12        # move stack pointer back up to what it was  
    jr   $ra               # return from sum_product
```

source code for more_calls.s

23

24

Example - more complex Calls - MIPS (product)

- a function which doesn't call other functions is called a leaf function
- its code can be simpler

```
int product(int x, int y) {  
    return x * y;  
}
```

source code for more_calls.c

```
product:          # product doesn't call other functions  
                  # so it doesn't need to save any registers  
    mul  $v0, $a0, $a1 # return argument * argument 2  
    jr   $ra             #
```

source code for more_calls.s

Example - strlen using array - C

C

```
int main(void) {  
    int i = my_strlen("Hello");  
    printf("%d\n", i);  
    return 0;  
}  
  
int my_strlen(char *s) {  
    int length = 0;  
    while (s[length] != 0) {  
        length++;  
    }  
    return length;  
}
```

source code for strlen_array.c

Simple C

```
int main(void) {  
    int i = my_strlen("Hello");  
    printf("%d\n", i);  
    return 0;  
}  
  
int my_strlen(char *s) {  
    int length = 0;  
    loop:  
        if (s[length] == 0) goto end;  
        length++;  
    goto loop;  
end:  
    return length;  
}
```

source code for strlen_array.simple.c

25

26

Example - pointer - C

```
int main(void) {  
    int i;  
    int *p;  
    p = &answer;  
    i = *p;  
    printf("%d\n", i); // prints 42  
    *p = 27;  
    printf("%d\n", answer); // prints 27  
    return 0;  
}
```

source code for pointer.c

Example - pointer - MIPS

main:

```
la   $t0, answer      # p = &answer;  
lw   $t1, ($t0)        # i = *p;  
move $a0, $t1           # printf("%d\n", i);  
li   $v0, 1  
syscall  
li   $a0, '\n'          # printf("%c", '\n');  
li   $v0, 11  
syscall  
li   $t2, 27            # *p = 27;  
sw   $t2, ($t0)  
lw   $a0, answer         # printf("%d\n", answer);  
li   $v0, 1  
syscall  
li   $a0, '\n'          # printf("%c", '\n');  
li   $v0, 11  
syscall
```

27

28

```

int main(void) {
    int i = my_strlen("Hello");
    printf("%d\n", i);
    return 0;
}

int my_strlen(char *s) {
    int length = 0;
    while (s[length] != 0) {
        length++;
    }
    return length;
}

```

source code for strlen_array.c

```

la    $a0, string      # my_strlen("Hello");
jal   my_strlen
move $a0, $v0           # printf("%d", i);
li    $v0, 1
syscall
li    $a0, '\n'         # printf("%c", '\n');
li    $v0, 11
syscall
lw    $ra, 0($sp)       # recover $ra from $stack
addi $sp, $sp, 4         # move stack pointer back up to what it was when
li    $v0, 0             # return 0 from function main
jr    $ra                 #

```

source code for strlen_array.s

29

30

Storing A Local Variables On the Stack

- some local (function) variables must be stored on stack
- e.g. variables such as arrays and structs

```

int main(void) {           main:
    int squares[10];          sub  $sp, $sp, 40
    int i = 0;                li   $t0, 0
    while (i < 10) {          loop0:
        squares[i] = i * i;    mul  $t1, $t0, 4
        i++;                  add   $t2, $t1, $sp
    }                         mul   $t3, $t0, $t0
                                sw    $t3, ($t2)
                                add   $t0, $t0, 1
                                b     loop0
end0:

```

source code for squares.c

source code for squares.s

Example - strlen using pointer - C

```

int main(void) {
    int i = my_strlen("Hello");
    printf("%d\n", i);
    return 0;
}

int my_strlen(char *s) {
    int length = 0;
    while (s[length] != 0) {
        length++;
    }
    return length;
}

```

source code for strlen_array.c

31

32

What is a Frame Pointer

Example of Growing Stack Breaking Function Return

- frame pointer \$fp is a second register pointing to stack
- by convention set to point at start of stack frame
- provides a fixed point during function code execution
- useful for functions which grow stack (change \$sp) during execution
- makes it easier for debuggers to forensically analyze stack
- e.g if you want to print stack backtrace after error
- frame pointer is optional (in COMP1521 and generally)
- often omitted when fast execution or small code a priority

```
void f(int a) {  
    int length;  
    scanf("%d", &length);  
    int array[length];  
    // ... more code ...  
    printf("%d\n", a);  
}  
  
source code for frame_pointer.c
```

```
f:  
    sub  $sp, $sp, 4  
    sw   $ra, 0($sp)  
    li   $v0, 5  
    syscall  
    # allocate space for  
    # array on stack  
    mul  $t0, $v0, 4  
    sub  $sp, $sp, $t0  
    # ... more code ...  
    # breaks because $sp  
    # has changed  
    lw   $ra, 0($sp)  
    add  $sp, $sp, 4  
    jr   $ra
```

source code for frame_pointer.broken.s

33

34

Example of Frame Pointer Use

```
void f(int a) {  
    int length;  
    scanf("%d", &length);  
    int array[length];  
    // ... more code ...  
    printf("%d\n", a);  
}  
  
source code for frame_pointer.c
```

```
f:  
    sub  $sp, $sp, 8  
    sw   $fp, 4($sp)  
    sw   $ra, 0($sp)  
    add  $fp, $sp, 8  
    li   $v0, 5  
    syscall  
    mul  $t0, $v0, 4  
    sub  $sp, $sp, $t0  
    # ... more code ...  
    lw   $ra, -4($fp)  
    move $sp, $fp  
    lw   $fp, 0($fp)  
    jr   $ra
```

source code for frame_pointer.s

35

- Operating system sits between the user and the hardware
- Operating system effectively provides a virtual machine to each user
- This virtual machine is much simpler and more convenient than real machine
- The virtual machine interface can be consistent across different hardware.
 - program can portably access hardware across different hardware configurations
 - linux available for almost all suitable hardware
- can coordinate/share access to resources between users
- can provide privileges/security

- needs hardware to provide a **privileged** mode which:
 - allows access to all hardware/memory
 - Operating System (kernel) runs in **privileged** mode
 - allows transfer to running code a **non-privileged** mode
- needs hardware to provide a **non-privileged** mode which:
 - prevents access to hardware
 - limits access to memory
 - provides mechanism to make requests to operating system
- operating system requests are called **system calls**
 - system calls transfers execution back to kernel code in **privileged** mode

1

2

System Call - What is It

- system call allow programs to request hardware operations
- system call transfers execution to OS code in **privileged** mode
 - includes arguments specifying details of request being made
 - OS checks operation is valid & permitted
 - OS carries out operation
 - transfers execution back to user code in **non-privileged** mode
- different operating system have different system calls
 - e.g Linux provides completely different system calls to Windows
- Linux provides 400+ system calls
- Operations likely to be provided by system calls:
 - read/write bytes to a file
 - request more memory
 - create a process (run a program)
 - terminate a process
 - send or receive information via a network

System Call in SPIM

- SPIM provides a virtual machine which can execute MIPS programs
- SPIM also provides a tiny operating system
- small number of SPIM system calls for I/O and memory allocation
- access is via the **syscall** instruction
- MIPS programs running on real hardware also use **syscall**
 - on linux **syscall**, will pass execution to linux kernel
- SPIM system calls are designed for students writing tiny programs
 - e.g. SPIM system call 1 - print an integer
- system calls on real operating systems more general
 - e.g. system call might be write n bytes
- in real operating system library functions like **printf** provide convenient operations
 - library functions like **printf** provide convenient operations

3

4

- **file systems** manage persistent stored data e.g. on magnetic disk or SSD
- On Unix-like systems:
 - a **file** is sequence (array) of zero or more bytes.
 - no meaning for bytes associated with file
 - file metadata doesn't record that it is e.g. ASCII, MP4, JPG, ...
 - Unix-like files are just bytes
 - a **directory** is an object containing zero or more files or directories.
- file systems maintain metadata for files & directories, e.g. permissions
- system calls provide operations to manipulate files.
- libc provides a low-level API to manipulate files.
- stdio.h provides more portable, higher-level API to manipulate files.

```
// hello world implemented with a direct syscall
#include <unistd.h>
int main(void) {
    char bytes[16] = "Hello, Andrew!\n";
    // argument 1 to syscall is system call number, 1 == write
    // remaining arguments are specific to each system call
    // write system call takes 3 arguments:
    // 1) file descriptor, 1 == stdout
    // 2) memory address of first byte to write
    // 3) number of bytes to write
    syscall(1, 1, bytes, 15); // prints Hello, Andrew! on stdout
    return 0;
}
```

source code for hello_syscalls.c

5

6

Using read & write system calls to copy stdin to stdout

```
// copy stdin to stdout with read & write syscalls
while (1) {
    char bytes[4096];
    // system call number 0 == read
    // read system call takes 3 arguments:
    // 1) file descriptor, 1 == stdin
    // 2) memory address to put bytes read
    // 3) maximum number of bytes read
    // returns number of bytes actually read
    long bytes_read = syscall(0, 0, bytes, 4096);
    if (bytes_read <= 0) {
        break;
    }
    syscall(1, 1, bytes, bytes_read); // prints bytes to stdout
}
```

source code for cat_syscalls.c

Using open system call to copy a file

```
// cp <file1> <file2> with syscalls, no error handling!
// system call number 2 == open, takes 3 arguments:
// 1) address of zero-terminated string containing file pathname
// 2) bitmap indicating whether to write, read, ... file
//     0x41 == write to file creating if necessary
// 3) permissions if file will be newly created
//     0644 == readable to everyone, writeable by owner
long read_file_descriptor = syscall(2, argv[1], 0, 0);
long write_file_descriptor = syscall(2, argv[2], 0x41, 0644);
while (1) {
    char bytes[4096];
    long bytes_read = syscall(0, read_file_descriptor, bytes, 4096);
    if (bytes_read <= 0) {
        break;
    }
    syscall(1, write_file_descriptor, bytes, bytes_read);
}
```

7

8

- Unix-like filenames are sequences of 1 or more bytes.
 - filenames can contain any byte except 0x00 and 0x2F
 - 0x00 bytes (ASCII '\0') used to terminate filenames
 - 0x2F bytes (ASCII '/') used to separate components of pathnames.
 - maximum filename length, depends on file system, typically 255
- Two filenames can not be used - they have a special meaning:
 - . current directory
 - .. parent directory
- Some programs (shell, ls) treat filenames starting with . specially.

9

- Files & directories accessed via pathnames, e.g:
/home/z5555555/lab07/main.c
- *absolute* pathnames start with a leading / and give full path from root
e.g. /usr/include/stdio.h, /cs1521/public_html/
- every process (running process) has an associated *absolute* pathname called the *current working directory* (CWD)
- shell command pwd prints CWD
- *relative* pathname do not start with a leading / e.g.
../../another/path/prog.c, ./a.out, main.c
- *relative* pathnames appended to CWD of process using them
- Assume process CWD is /home/z5555555/lab07/
main.c translated to absolute path /home/z5555555/lab07/main.c
./a.out translated to absolute path /home/z5555555/./a.out
which is equivalent to absolute path /home/z5555555/a.out

10

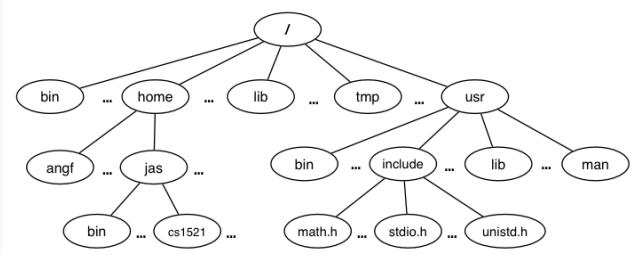
Everything is a File

- Originally files only managed data stored on a magnetic disk.
- Unix philosophy is: *Everything is a File*.
- File system can be used to access:
 - files
 - directories (folders)
 - storage devices (disks, SSD, ...)
 - peripherals (keyboard, mouse, USB, ...)
 - system information
 - inter-process communication
 - ...

9

Unix/Linux File System

Unix/Linux file system is tree-like



- We think of file-system as a *tree*
- But beware if you follow symbolic links it is a *graph*.
 - and you may infinitely loop attempting to traverse a file system
 - but only if you follow symbolic links

Metadata for file system objects is stored in **inodes**, which hold

- location of file contents in file systems
- file type (regular file, directory, ...)
- file size in byte
- file ownership
- file access permissions - who can read, write, execute the file
- timestamps - time of creation/access/update

Note: file systems add much complexity to improve performance

- e.g. very small files might be stored in an inode itself

13

14

File Access: Behind the Scenes

Access to files by name proceeds (roughly) as...

- open directory and scan for *name*
- if not found, "No such file or directory"
- if found as (*name,inumber*), access inode table *inodes[inumber]*
- collect file metadata and...
 - check file access permissions given current user/group
 - if don't have required access, "Permission denied"
 - collect information about file's location and size
 - update access timestamp
- use data in inode to access file contents

- unix-like file systems effectively have an array of inodes
- every inode has a ***inode-number*** (or *i-number*)- its index in this array
- directories are effectively a list of (name, inode-number) pairs
- inode-number uniquely identify files within a filesystem
 - just a uid uniquely identifies a student within UNSW
- **ls -i** prints *inode-number*, e.g.:

```
$ ls -i file.c
109988273 file.c
$
```

Hard Links & Symbolic Links

File system *links* allow multiple paths to access the same file

Hard links

- multiple names referencing the same file (inode)
- the two entries must be on the same filesystem
- all hard links to a file have equal status
- file destroyed when last hard link removed
- can not create a (extra) hard link to directories

Symbolic links (symlinks)

- point to another path name
- accessing the symlink (by default) accesses the file being pointed to
- symbolic link can point to a directory
- symbolic link can point to a pathname on another filesystems

15

16

```
$ echo 'Hello Andrew' >hello
$ ln hello hola      # create hard link
$ ln -s hello selamat # create symbolic link
$ ls -l hello hola selamat
-rw-r--r-- 2 andrewt 13 Oct 23 16:18 hello
-rw-r--r-- 2 andrewt 13 Oct 23 16:18 hola
lrwxrwxrwx 1 andrewt  5 Oct 23 16:20 selamat -> hello
$ cat hello
Hello Andrew
$ cat hola
Hello Andrew
$ cat selamat
Hello Andrew
```

Unix presents a uniform interface to file system objects

- functions/syscalls manipulate objects as a *stream of bytes*
- accessed via a *file descriptor*
 - file descriptors are small integers
 - index to a per-process operating system table (array)

Some common operations:

- **open()** — open a file system object, returning a file descriptor
- **close()** — stop using a file descriptor
- **read()** — read some bytes into a buffer from a file descriptor
- **write()** — write some bytes from a buffer to a file descriptor
- **lseek()** — move to a specified offset within a file
- **stat()** — get file system object metadata

17

18

libc Types for File System Operations

Unix defines a range of file-system-related types:

- **off_t** — offsets within files
 - typically `int64_t` - signed to allow backward references
- **size_t** — number of bytes in some object
 - typically `uint64_t` - unsigned since objects can't have negative size
- **ssize_t** — sizes of read/written bytes
 - like `*size_t`, but signed to allow for error values
- **struct stat** — file system object metadata
 - stores information *about* file, not its contents
 - requires other types: `ino_t`, `dev_t`, `time_t`, `uid_t`, ...

open - libc

`int open(char *pathname, int flags)`

- open file at `pathname`, according to `flags`
- `flags` is a bit-mask defined in `<fcntl.h>`
 - `O_RDONLY` — open for reading
 - `O_WRONLY` — open for writing
 - `O_APPEND` — append on each write
 - `O_RDWR` — open object for reading and writing
 - `O_CREAT` — create file if doesn't exist
 - `O_TRUNC` — truncate to size 0
- `flags` can be combined e.g. `(O_WRONLY|O_CREAT)`
- if successful, return file descriptor (small non-negative `int`)
- if unsuccessful, return `-1` and set `errno`

19

20

close - libc

```
int close(int fd)
▪ release open file descriptor fd
▪ if successful, return 0
▪ if unsuccessful, return -1 and set errno
    ▪ could be unsuccessful if fd is not an open file descriptor
        e.g. if fd has already been closed
```

An aside: removing a file e.g. via rm

- removes the file's entry from a directory
- but the inode and data persist until
 - all references to the inode from other directories are removed
 - all processes accessing the file close() their file descriptor
- after this, the inode and the space used for file contents is recycled

read

```
ssize_t read(int fd, void *buf, size_t count)
▪ read (up to) count bytes from fd into buf
    ▪ buf should point to array of at least count bytes
    ▪ read does (can) not check buf points to enough space
    ▪ if successful, number of bytes actually read is returned
    ▪ 0 returned, if no more bytes to read
    ▪ -1 returned if error and errno set to reason
    ▪ next call to read will return next bytes from file
    ▪ repeated calls to reads will yield entire contents of file
        ▪ associated with a file descriptor is "current position" in file
        ▪ can also modify this position with lseek
```

21

22

write - libc

```
ssize_t write(int fd, const void *buf, size_t count)
▪ attempt to write count bytes from buf into
    stream identified by file descriptor fd
▪ if successful, number of bytes actually written is returned
▪ if unsuccessful, return -1 and set errno
▪ does (can) not check buf points to count bytes of data
▪ next call to write will follow bytes already written
▪ file often created by repeated calls to write
    ▪ associated with a file descriptor is "current position" in file
    ▪ can also modify this position with lseek
```

Hello libc

```
// hello world implemented with libc
#include <unistd.h>
int main(void) {
    char bytes[16] = "Hello, Andrew!\n";
    // write takes 3 arguments:
    // 1) file descriptor, 1 == stdout
    // 2) memory address of first byte to write
    // 3) number of bytes to write
    write(1, bytes, 15); // prints Hello, Andrew! on stdout
    return 0;
}
```

source code for hello_libc.c

23

24

Using read & write to copy stdin to stdout

```
while (1) {
    char bytes[4096];
    // system call number 0 == read
    // read system call takes 3 arguments:
    //   1) file descriptor, 1 == stdin
    //   2) memory address to put bytes read
    //   3) maximum number of bytes read
    // returns number of bytes actually read
    ssize_t bytes_read = read(0, bytes, 4096);
    if (bytes_read <= 0) {
        break;
    }
    write(1, bytes, bytes_read); // prints bytes to stdout
}
```

source code for cat libc.c

Using open to copy a file

```
// open takes 3 arguments:
//   1) address of zero-terminated string containing pathname of file
//   2) bitmap indicating whether to write, read, ... file
//   3) permissions if file will be newly created
//       0644 == readable to everyone, writeable by owner
int read_file_descriptor = open(argv[1], O_RDONLY);
int write_file_descriptor = open(argv[2], O_WRONLY | O_CREAT, 0644);
while (1) {
    char bytes[4096];
    ssize_t bytes_read = read(read_file_descriptor, bytes, 4096);
    if (bytes_read <= 0) {
        break;
    }
    write(write_file_descriptor, bytes, bytes_read);
}
```

source code for cp libc.c

25

26

lseek

```
off_t lseek(int fd, off_t offset, int whence)


- change the 'current position' in the file of fd
- offset is in units of bytes, and can be negative
- whence can be one of ...
  - SEEK_SET — set file position to Offset from start of file
  - SEEK_CUR — set file position to Offset from current position
  - SEEK_END — set file position to Offset from end of file
- seeking beyond end of file leaves a gap which reads as 0's
- seeking back beyond start of file sets position to start of file

```

Example: lseek(fd, 0, SEEK_END); (move to end of file)

Using lseek to read the last byte then the first byte of a file

```
int read_file_descriptor = open(argv[1], O_RDONLY);
char bytes[1];
// move to a position 1 byte from end of file
// then read 1 byte
lseek(read_file_descriptor, -1, SEEK_END);
read(read_file_descriptor, bytes, 1);
printf("The last byte of the file is 0x%02x\n", bytes[0]);
// move to a position 0 bytes from start of file
// then read 1 byte
lseek(read_file_descriptor, 0, SEEK_SET);
read(read_file_descriptor, bytes, 1);
printf("The first byte of the file is 0x%02x\n", bytes[0]);
```

source code for lseek.c

27

28

Using lseek to read bytes in the middle of a file

```
printf("The first byte of the file is 0x%02x\n", bytes[0]);
// move to a position 41 bytes from start of file
// then read 1 byte
lseek(read_file_descriptor, 41, SEEK_SET);
read(read_file_descriptor, bytes, 1);
printf("The 42nd byte of the file is 0x%02x\n", bytes[0]);
// move to a position 58 bytes from current position
// then read 1 byte
lseek(read_file_descriptor, 58, SEEK_CUR);
read(read_file_descriptor, bytes, 1);
printf("The 100th byte of the file is 0x%02x\n", bytes[0]);
```

source code for lseek.c

stat

```
int stat(const char *pathname, struct stat *statbuf)
■ returns metadata associated with pathname in statbuf
■ metadata returned includes:
    ■ inode number
    ■ type (file, directory, symbolic link, device)
    ■ size of file in bytes (if it is a file)
    ■ permissions (read, write, execute)
    ■ times of last access/modification/status-change
■ returns -1 and sets errno if metadata not accessible
int fstat(int fd, struct stat *statbuf)
■ same as stat() but gets data via an open file descriptor
int lstat(const char *pathname, struct stat *statbuf)
■ same as stat() but doesn't follow symbolic links
```

29

30

definition of struct stat

```
struct stat {
    dev_t      st_dev;          /* ID of device containing file */
    ino_t      st_ino;          /* Inode number */
    mode_t     st_mode;         /* File type and mode */
    nlink_t    st_nlink;        /* Number of hard links */
    uid_t      st_uid;          /* User ID of owner */
    gid_t      st_gid;          /* Group ID of owner */
    dev_t      st_rdev;         /* Device ID (if special file) */
    off_t      st_size;         /* Total size, in bytes */
    blksize_t  st_blksize;      /* Block size for filesystem I/O */
    blkcnt_t   st_blocks;       /* Number of 512B blocks allocated */
    struct timespec st_atim;   /* Time of last access */
    struct timespec st_mtim;   /* Time of last modification */
    struct timespec st_ctim;   /* Time of last status change */
};
```

st_mode mode field of struct stat

st_mode is a bitwise-or of these values (& others):

S_IFLNK	0120000	symbolic link
S_IFREG	0100000	regular file
S_IFBLK	0060000	block device
S_IFDIR	0040000	directory
S_IFCHR	0020000	character device
S_IFIFO	0010000	FIFO
S_IRUSR	0000400	owner has read permission
S_IWUSR	0000200	owner has write permission
S_IXUSR	0000100	owner has execute permission
S_IRGRP	0000040	group has read permission
S_IWGRP	0000020	group has write permission
S_IXGRP	0000010	group has execute permission
S_IROTH	0000004	others have read permission
S_IWOTH	0000002	others have write permission
S_IXOTH	0000001	others have execute permission

31

32

Using stat

```
struct stat s;
if (stat(pathname, &s) != 0) {
    perror(pathname);
    exit(1);
}
printf("s.st_ino = %10ld # Inode number\n", s.st_ino);
printf("s.st_mode = %10o # File mode \n", s.st_mode);
printf("s.st_nlink =%10ld # Link count \n", (long)s.st_nlink);
printf("s.st_uid = %10u # Owner uid\n", s.st_uid);
printf("s.st_gid = %10u # Group gid\n", s.st_gid);
printf("s.st_size = %10ld # File size (bytes)\n", (long)s.st_size);
printf("s.st_mtime =%10ld # Modification time (seconds since 1/1/70\n", s.st_mtime);
```

source code for stat.c

mkdir

- ```
int mkdir(const char *pathname, mode_t mode)
```
- create a new directory called **pathname** with permissions **mode**
  - if **pathname** is e.g. a/b/c/d
    - all of the directories a, b and c must exist
    - directory d must be writeable to the caller
    - directory d must not already exist
  - the new directory contains two initial entries
    - . is a reference to itself
    - .. is a reference to its parent directory
  - returns 0 if successful, returns -1 and sets errno otherwise

Example:

```
mkdir("newDir", 0755);
```

33

34

## Other useful Linux (POSIX) functions

```
chdir(char *path) // change current working directory

getcwd(char *buf, size_t size) // get current working directory

rename(char *oldpath, char *newpath) // rename a file/directory

link(char *oldpath, char *newpath) // create hard link to a file

symlink(char *target, char *linkpath) // create a symbolic link

unlink(char *pathname) // remove a file/directory...

chmod(char *pathname, mode_t mode) // change permission of file/...
```

## stdio.h

stdio.h functions more portable more convenient than open/read/write/... use them by default

- stdio.h equivalent to open is **fopen**

```
FILE *fopen(const char *pathname, const char *mode)
```

- **mode** is string of 1 or more characters including:
    - r open text file for reading.
    - w open text file for writing truncated to 0 zero length if it exists created if does not exist
    - a open text file for writing writes append to it if it exists created if does not exist
  - fopen returns a **FILE \*** pointer
  - **FILE** is an opaque struct - we can not access fields
- ```
int fclose(FILE *stream)
```
- stdio.h equivalent to close

35

36

```

int fgetc(FILE *stream)           // read a byte
int fputc(int c, FILE *stream)    // write a byte

char *fputs(char *s, FILE *stream) // write a string

char *fgets(char *s, int size, FILE *stream) // read a line

// formatted input
int fscanf(FILE *stream, const char *format, ...)

// formatted output
int fprintf(FILE *stream, const char *format, ...)

// read array of bytes
size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream);
// write array of bytes
size_t fwrite(const void *ptr, size_t size, size_t nmemb,
             FILE *stream)

```

37

stdio.h - other operations on

```

int fseek(FILE *stream, long offset, int whence);


- fseek is stdio equivalent to lseek
- like lseek offset can be positive or negative
- like lseek whence can be SEEK_SET, SEEK_CUR or SEEK_END making
      offset relative to file start, current position or file end



int fflush(FILE *stream);


- flush any buffered data on writing stream

```

```

int fclose(FILE *stream)


- equivalent to close


@hello_fputc.c @hello_fputs.c @hello_fwrite.c @hello_fprintf.c @cat_fgetc.c
@cat_fgets.c @cat_fwrite.c @cp_fgetc.c @cp_fwrite.c @create_file_fopen.c
@create_append_truncate_fopen.c

```

As we often read/write to stdin/stdout stdio.h provides convenience functions which only read/write stdin/stdout

```

int getchar()           // fgetc(stdin)
int putchar(int c)     // fputc(c, stdin)

int puts(char *s)       // fputs(s, stdout)

int scanf(char *format, ...) // fscanf(stdin, format, ...)
int printf(char *format, ...) // fprintf(stdout, format, ...)

char *gets(char *s); // NEVER USE

```

38

stdio.h - I/O to strings

stdio.h provides useful functions which operate on strings

```

int snprintf(char *str, size_t size, const char *format, ...);


- like printf, but output goes to char array str
- handy for creating strings passed to other functions
- do not use unsafe related function: 'sprintf'



int sscanf(const char *str, const char *format, ...);


- like scanf, but input comes from char array str



int sprintf(char *str, const char *format, ...); // DO NOT USE


- like snprintf but dangerous because can overflow str


@myio_unbuffered.c @myio_input_buffered.c @myio_output_buffered.c

```

39

40

Operating systems provide a *file system*

- as an abstraction over physical storage devices (e.g. disks)
- providing named access to chunks of related data (files)
- providing access (sequential/random) to the contents of files
- allowing files to be arranged in a hierarchy of directories
- providing control over access to files and directories
- managing other metadata associated with files (size, location, ...)

Operating systems also manage other resources

- memory, processes, processor time, i/o devices, networking, ...

Huge number of character representations (encodings) exist
you need know only two:

- ASCII (ISO 646)
 - single byte values, only low 7-bit used, top bit always 0
 - can encode roman alphabet a-zA-Z, digits 0-9 , punctuation, control chars
 - complete alphabet for English, Bahasa
 - no diacritics, e.g: ç , so missing a little of alphabet for other latin languages, e.g.: German, French, Spanish, Italian, Swedish, Tagalog, Swahili
 - characters for most of world's languages completely missing
- UTF-8 (Unicode)
 - contains all ASCII (single-byte) values
 - also has 2-4 byte values, top bit always 1 for bytes of multi-byte values
 - contains symbols for essentially all human languages plus other symbols, e.g.:



- Uses values in the range 0x00 to 0x7F (0..127)
- Characters partitioned into sequential groups
 - control characters (0..31) ... e.g. '\n'
 - punctuation chars (32..47,91..96,123..126)
 - digits (48..57) ... '0'..'9'
 - upper case alphabetic (65..90) ... 'A'..'Z'
 - lower case alphabetic (97..122) ... 'a'..'z'
- Sequential nature of groups allow ordination e.g.
'3' - '0' == 3 'J' - 'A' == 10
- See [man 7 ascii](#)

1

2

Unicode

UTF-8 Encoding

#bytes	#bits	Byte 1	Byte 2	Byte 3	Byte 4
1	7	0xxxxxx	-	-	-
2	11	110xxxx	10xxxxx	-	-
3	16	1110xxx	10xxxxx	10xxxxx	-
4	21	11110xx	10xxxxx	10xxxxx	10xxxxx

- The 127 1-byte codes are compatible with ASCII
- The 2048 2-byte codes include most Latin-script alphabets
- The 65536 3-byte codes include most Asian languages
- The 2097152 4-byte codes include symbols and emojis and ...

ch	code-point	unicode binary	UTF-8 encoding
\$	U+0024	0100100	00100100
¢	U+00A2	00010100010	11000010 10100010
€	U+20AC	0010000010101100	11100010 10000010 10101100

3

4

```

printf("The unicode code point U+1F600 encodes in UTF-8\n");
printf("as 4 bytes: 0xF0 0x9F 0x98 0x80\n");
printf("We can output the 4 bytes like this: \xF0\x9F\x98\x80\n");
printf("Or like this: ");
putchar(0xF0);
putchar(0x9F);
putchar(0x98);
putchar(0x80);
putchar('\n');

```

source code for hello.c

```

uint8_t encoding[5] = {0};
if (code_point < 0x80) {
    encoding[0] = code_point;
} else if (code_point < 0x800) {
    encoding[0] = 0xC0 | (code_point >> 6);
    encoding[1] = 0x80 | (code_point & 0x3f);
} else if (code_point < 0x10000) {
    encoding[0] = 0xE0 | (code_point >> 12);
    encoding[1] = 0x80 | ((code_point >> 6) & 0x3f);
    encoding[2] = 0x80 | (code_point & 0x3f);
} else if (code_point < 0x200000) {
    encoding[0] = 0xF0 | (code_point >> 18);
    encoding[1] = 0x80 | ((code_point >> 12) & 0x3f);
    encoding[2] = 0x80 | ((code_point >> 6) & 0x3f);
    encoding[3] = 0x80 | (code_point & 0x3f);
}

```

5

source code for utf8_encode.c

6

Converting Unicode Codepoints to UTF-8

```

printf("U+%x  UTF-8: ", code_point);
for (uint8_t *s = encoding; *s != 0; s++) {
    printf("0x%02x ", *s);
}
printf("%s\n", encoding);
}

int main(void) {
    print_utf8_encoding(0x42);
    print_utf8_encoding(0x00A2);
    print_utf8_encoding(0x10be);
    print_utf8_encoding(0x1F600);
}

```

source code for utf8_encode.c

Summary of UTF-8 Properties

- Compact, but not minimal encoding; encoding allows you to resync immediately if bytes lost from a stream.
- ASCII is a subset of UTF-8 - complete backwards compatibility!
- All other UTF-8 bytes > 127 (0x7f)
 - no byte of multi-byte UTF-8 encoding is valid ASCII.
- No byte of multi-byte UTF-8 encoding is 0
 - can still use store UTF-8 in null-terminated strings.
- 0x2F (ASCII '/') and 0x00 can not appear in multi-byte characters
 - hence can use UTF-8 for Linux/Unix filenames
- C programs can treat UTF-8 similarly to ASCII.
- Beware: number of bytes in UTF-8 string != number of characters.

```
virtual_memory # Processes
```

A process is an instance of an executing program

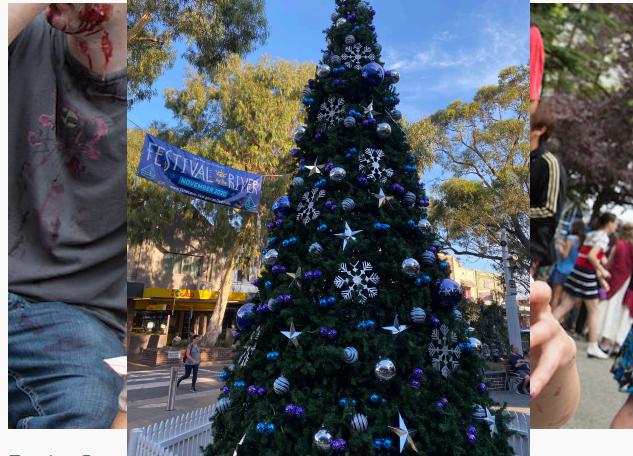
Each process has an *execution state*, defined by

- current values of CPU registers
- current contents of its (virtual) memory
- information about open files, sockets, etc.

On Unix/Linux:

- each process had a unique process ID (pid)
- positive integer - type `pid_t` defined in `unistd.h`
- process 0 is effectively part of the operating system
- process 1 (**`init`**) - used to boot the system
- some parts of operating system may run as processes
- low-numbered processes are typically system-related process started at boot-time

Aside: Zombie Processes



Zombie Process

Photo credit: Kenny Louie, Flickr.com

Process Parents

Each process has a *parent process*

- initially it is the process that created it
- if a process' parent terminates, its parent becomes process 1

Unix provides a range of commands for manipulating processes, e.g.:

- `sh ...` for creating processes via object-file name
- `ps ...` show process information
- `w ...` show per-user process information
- `top ...` show high-cpu-usage process information
- `kill ...` send a signal to a process

1

2

Aside: zombie Processes

- a process can't terminate until its parent is notified
- if `exit()` called, operating system sends `SIGCHLD` signal to parent
- `exit()` will not return until parent handles `SIGCHLD`
- *Zombie process* = exiting process waiting for parent to handle `SIGCHLD`
- all processes become zombies until `SIGCHLD` handled
- bug in parent that ignores `SIGCHLD` creates long-term zombie processes
 - wastes some operating system resources
- *Orphan process* = a process whose parent has exited
 - when parent exits, orphan is assigned pid=1 (`init`) as its parent
 - `init` should always handle `SIGCHLD` when process exits

3

4

Multi-Tasking

On a typical modern operating system

- multiple processes are active "simultaneously" (*multi-tasking*)
- operating systems provides a virtual machine to each process
 - each process executes as if the only process running on the machine
 - e.g. each process has its own address space (N bytes, addressed 0..N-1)

When there are multiple processes running on the machine

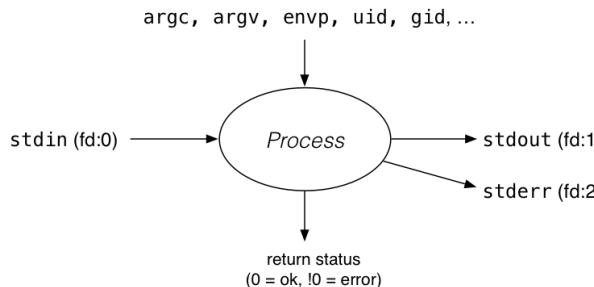
- each process uses the CPU until *pre-empted* or exits
- then another process uses the CPU until it too is pre-empted
- eventually, the first process will get another run on the CPU



Overall impression: three programs running simultaneously

Unix/Linux Processes

Environment for processes running on Unix/Linux systems



Processes

What can cause a process to be pre-empted?

- it runs "long enough" and the OS replaces it by a waiting process
- it needs to wait for input, or output or ...

On pre-emption ..

- the process's entire state must be saved
- the new process's state must be restored
- this change is called a **context switch**
- **context switches** are expensive

The operating system's process scheduling attempts to:

- fairly sharing the CPU(s) among competing processes
- minimize response delays (lagginess) for interactive users
- meet other real-time requirements (e.g. self-driving car)
- minimize number of expensive context switches

5

6

Process-related Unix/Linux Functions/System Calls

- `posix_spawn()` ... create a new process, see also
 - `clone()` ... duplicate current process
 - address space can be shared to implement threads
 - only use clone if posix_spawn can't do what you want
- `fork()` ... duplicate current process - do not use in new code
- `execvp()` ... replace current process
- `system()` `popen()` ... create a new process via a shell (unsafe)
- `exit()` ... terminate current process, see also
 - `_exit()` ... terminate current process immediately
 - atexit functions not be called: stdio buffers not flushed
- `getpid()` ... get process ID
- `getpgid()` ... get process group ID
- `waitpid()` ... wait for state change in child process

7

8

posix_spawn() - run a new process

```
#include <spawn.h>

int posix_spawn(pid_t *pid, const char *path,
               const posix_spawn_file_actions_t *file_actions,
               const posix_spawnattr_t *attrp,
               char *const argv[], char *const envp[]);
```

- creates new process, running program at path
- argv specifies argv of new program
- envp specifies environment of new program
- *pid set to process id of new program
- file_actions specifies file actions to be performed before running program
 - can be used to re-direct stdin or stdout to file or pipe
 - advanced topic
- attrp specifies attributes for new process

9

Simple example using posix_spawn() to run /bin/date

```
pid_t pid;
extern char **environ;
char *date_argv[] = {"/bin/date", "--utc", NULL};
// spawn "/bin/date" as a separate process
if (posix_spawn(&pid, "/bin/date", NULL, NULL, date_argv, environ)
    perror("spawn");
    exit(1);
}
// wait for spawned processes to finish
int exit_status;
if (waitpid(pid, &exit_status, 0) == -1) {
    perror("waitpid");
    exit(1);
}
printf("/bin/date exit status was %d\n", exit_status);
```

source code for spawn.c

10

fork() - clone yourself

```
#include <sys/types.h>
#include <unistd.h>
```

```
pid_t fork(void);



- creates new process by duplicating the calling process
- new process is the child, calling process is the parent
- child has a different process ID (pid) to the parent
- in the child, fork() returns 0
- in the parent, fork() returns the pid of the child
- if the system call fails, fork() returns -1
- child inherits copies of parent's address space and open file descriptors
- do not use in new code use posix_spawn instead
  - fork appears simple but prone to subtle bugs

```

11

Simple example of using fork()

```
// fork creates 2 identical copies of program
// only return value is different
pid_t pid = fork();
if (pid == -1) {
    perror("fork"); // print why the fork failed
} else if (pid == 0) {
    printf("I am the child because fork() returned %d.\n", pid);
} else {
    printf("I am the parent because fork() returned %d.\n", pid);
}
```

source code for fork.c

```
$ gcc fork.c
$ a.out
I am the parent because fork() returned 2884551.
I am the child because fork() returned 0.
```

12

```
#include <unistd.h>

int execvp(const char *file, char *const argv[]);



- replaces current process by executing file
  - file must be an executable: binary or script starting with #!
- argv specifies argv of new program
- most of the current process is reset
  - e.g. new virtual address space is created, signal handlers reset
- new process inherits open file descriptors from original process
- on error, returns -1 and sets errno
- if successful, does not return

```

13

```
char *echo_argv[] = {"/bin/echo", "good-bye", "cruel", "world", NULL};
execv("/bin/echo", echo_argv);
// if we get here there has been an error
perror("execv");
```

source code for exec.c

```
$ gcc exec.c
$ a.out
good-bye cruel world
$
```

Simple example using fork() and exec() to run /bin/date

```
pid_t pid = fork();
if (pid == -1) {
    perror("fork"); // print why fork failed
} else if (pid == 0) { // child
    char *date_argv[] = {"/bin/date", "--utc", NULL};
    execv("/bin/date", date_argv);
    perror("execve"); // print why exec failed
} else { // parent
    int exit_status;
    if (waitpid(pid, &exit_status, 0) == -1) {
        perror("waitpid");
        exit(1);
    }
    printf("/bin/date exit status was %d\n", exit_status);
}
```

source code for fork_exec.c

system() - convenient but unsafe way to run another program

```
#include <stdlib.h>
int system(const char *command);



- runs command via /bin/sh
- waits for command to finish and returns exit status
- convenient but brittle and highly vulnerable to security exploits
- use for quick debugging and throw-away programs only


// run date --utc to print current UTC
int exit_status = system("/bin/date --utc");
printf("/bin/date exit status was %d\n", exit_status);
return 0;
```

source code for system.c

15

14

16

running ls -ld via posix_spawn

```
char *ls_argv[argc + 2];
ls_argv[0] = "/bin/ls";
ls_argv[1] = "-ld";
for (int i = 1; i <= argc; i++) {
    ls_argv[i + 1] = argv[i];
}
pid_t pid;
extern char **environ;
if (posix_spawn(&pid, "/bin/ls", NULL, NULL, ls_argv, environ) != 0) {
    perror("spawn");
    exit(1);
}
```

source code for lsld_spawn.c

running ls -ld via posix_spawn

```
int exit_status;
if (waitpid(pid, &exit_status, 0) == -1) {
    perror("waitpid");
    exit(1);
}
// exit with whatever status ls exited with
return exit_status;
```

source code for lsld_spawn.c

17

18

running ls -ld via system

```
char *ls = "/bin/ls -ld";
int command_length = strlen(ls);
for (int i = 1; i < argc; i++) {
    command_length += strlen(argv[i]) + 1;
}
// create command as string
char command[command_length + 1];
strcpy(command, ls);
for (int i = 1; i <= argc; i++) {
    strcat(command, " ");
    strcat(command, argv[i]);
}
int exit_status = system(command);
```

source code for lsld_system.c

getpid & getppid

```
#include <sys/types.h>
#include <unistd.h>

pid_t getpid(void);
pid_t getppid(void);



- getpid returns the process ID of the current process
- getppid returns process ID of the the parent of current process

```

19

20

waitpid

```
#include <sys/types.h>
#include <sys/wait.h>

pid_t waitpid(pid_t pid, int *wstatus, int options);
pid_t wait(int *wstatus);



- waitpid pauses current process until process pid changes state
  - where state changes include finishing, stopping, re-starting, ...
- ensures that child resources are released on exit
- special values for pid ...
  - if pid = -1, wait on any child process
  - if pid = 0, wait on any child in process group
  - if pid > 0, wait on the specified process



pid_t wait(int *status)



- equivalent to waitpid(-1, &status, 0)
- pauses until one of the child processes terminates

```

linux/environment variables

- when linux/unix program are passed **environment variables**
- **environment variables** are array of strings of form `name=value`
- array is NULL-terminated
- access via global variable `environ`
- many C implementation also provide as 3rd parameter to `main`:


```
int main(int argc, char *argv[], char *env[])
```
- most program use `getenv` & `setenv` to access environment variables
- can access environment variables directly, eg:

```
// print all environment variables
extern char **environ;
for (int i = 0; environ[i] != NULL; i++) {
    printf("%s\n", environ[i]);
}
```

source code for environ.c

waitpid

More on `waitpid(pid, &status, options)`

- `status` is set to hold info about `pid`
 - e.g. exit status if `pid` terminated
 - macros allow precise determination of state change (e.g. `WIFEXITED(status)`, `WCOREDUMP(status)`)
- `options` provide variations in `waitpid()` behaviour
 - default: wait for child process to terminate
 - `WNOHANG`: return immediately if no child has exited
 - `WCONTINUED`: return if a stopped child has been restarted

For more information: `man 2 waitpid`

21

22

accessing an environment variable with getenv

```
#include <stdlib.h>
char *getenv(const char *name);



- search environment variable array for name=value
- returns value
- returns NULL if name not in environment variable array



// print value of environment variable STATUS
char *value = getenv("STATUS");
printf("Environment variable 'STATUS' has value '%s'\n", value);
```

source code for get_status.c

23

24

setting an environment variables with setenv

```
#include <stdlib.h>
int setenv(const char *name, const char *value, int overwrite);
    ▪ adds name=value to environment variable array
    ▪ if name in array, value changed if overwrite is non-zero

// set environment variable STATUS
setenv("STATUS", "great", 1);
char *getenv_argv[] = {"./get_status", NULL};
pid_t pid;
extern char **environ;
if (posix_spawn(&pid, "./get_status", NULL, NULL,
    getenv_argv, environ) != 0) {
    perror("spawn");
    exit(1);
}
```

source code for set_status.c

changing behaviour with an environment variable

```
pid_t pid;
char *date_argv[] = { "/bin/date", NULL };
char *date_environment[] = { "TZ=Australia/Perth", NULL };
// print time in Perth
if (posix_spawn(&pid, "/bin/date", NULL, NULL, date_argv,
    date_environment) != 0) {
    perror("spawn");
    return 1;
}
int exit_status;
if (waitpid(pid, &exit_status, 0) == -1) {
    perror("waitpid");
    return 1;
}
printf("/bin/date exit status was %d\n", exit_status);
```

source code for spawn_environment.c

exit() - terminate yourself

```
#include <stdlib.h>

void exit(int status);
    ▪ triggers any functions registered as atexit()
    ▪ flushes stdio buffers; closes open FILE *'s
    ▪ terminates current process
    ▪ a SIGCHLD signal is sent to parent
    ▪ returns status to parent (via waitpid())
    ▪ any child processes are inherited by init (pid 1)

Also void _exit(int status)
    ▪ terminates current process without triggering functions registered as atexit()
    ▪ stdio buffers not flushed
```

pipe() - stream bytes between processes

```
#include <unistd.h>
int pipe(int pipefd[2]);
```

- a pipe is a unidirectional byte stream provided by operating system
- pipefd[0] - set to file descriptor of read end of pipe
- pipefd[1] - set to file descriptor of write end of pipe
- bytes written to pipefd[1] will be read from pipefd[1]
- child processes (by default) inherit file descriptors including for pipe
- parent can send/receive bytes (not both) to child via pipe
- parent and child should both close the pipe file descriptor they are not using
 - e.g if bytes being written (sent) parent to child
 - parent should close read end pipefd[0]
 - child should close write end pipefd[1]
- pipe (and other) file descriptors can be used with stdio via fdopen

popen() - convenient but unsafe way to set up pipe

```
#include <stdio.h>
FILE *popen(const char *command, const char *type);
int pclose(FILE *stream);



- runs command via /bin/sh
- if type is "w" pipe to stdin of command created
- if type is "r" pipe from stdout of command created
- FILE * stream returned - get then use fgetc/fputc etc
- NULL returned if error
- close stream with pclose (not fclose)
  - pclose waits for command and returns exit status
- convenient but brittle and highly vulnerable to security exploits
- use for quick debugging and throw-away programs only

```

popen() - capturing output from a process

```
// popen passes string to a shell for evaluation
// brittle and highly-vulnerable to security exploits
// popen is suitable for quick debugging and throw-away programs or
FILE *p = popen("/bin/date --utc", "r");
if (p == NULL) {
    perror("");
    return 1;
}
char line[256];
if (fgets(line, sizeof line, p) == NULL) {
    fprintf(stderr, "no output from date\n");
    return 1;
}
printf("output captured from /bin/date was: '%s'\n", line);
pclose(p); // returns command exit status
```

source code for read_popen.c

29

30

popen() - sending input to a process

```
int main(void) {
    // popen passes command to a shell for evaluation
    // brittle and highly-vulnerable to security exploits
    // popen is suitable for quick debugging and throw-away programs
    //
    // tr a-z A-Z - passes stdin to stdout converting lower case to
FILE *p = popen("tr a-z A-Z", "w");
if (p == NULL) {
    perror("");
    return 1;
}
fprintf(p, "plz date me\n");
pclose(p); // returns command exit status
return 0;
}
```

source code for write_popen.c

posix_spawn and pipes (advanced topic)

```
int posix_spawn_file_actions_destroy(
    posix_spawn_file_actions_t *file_actions);
int posix_spawn_file_actions_init(
    posix_spawn_file_actions_t *file_actions);
int posix_spawn_file_actions_addclose(
    posix_spawn_file_actions_t *file_actions, int fildes);
int posix_spawn_file_actions_adddup2(
    posix_spawn_file_actions_t *file_actions, int fildes,
    int newfildes);



- functions to combine file operations with posix_spawn process creation
- awkward to understand & use - but robust
- example: capturing output from a process - source code for spawn\_read\_pipe.c
- example: sending input to a process - source code for spawn\_write\_pipe.c

```

31

32

- signal are simple form of interprocess-communication
- signals can be generated from a variety of sources
 - from another process via `kill()`
 - from the operating system (e.g. timer)
 - from within the process (e.g. system call)
 - from a fault in the process (e.g. div-by-zero)
- processes can define how they want to handle signals
 - using the `signal()` library function (simple)
 - using the `sigaction()` system call (powerful)
- signal `SIGKILL` always terminates receiving processes
- only owner of a processes can send signal to it

Default handling of signal can be:

- **Term** ... terminate the process
- **Ign** ... ignored - the signal does nothing
- **Core** ... terminate the process and dump memory image to file named core
- **Stop** ... pause the process
- **Cont** ... continue the process (if paused)

Processes can choose to ignore a signal.

Processes can set a custom *signal handler* for signal.

Except `SIGKILL` and `SIGSTOP` cannot be caught, blocked, or ignored.

See `man 7 signal` for details of signals and default handling.

1

2

Operating System-Generated Signals

Signals from internal process activity, e.g.

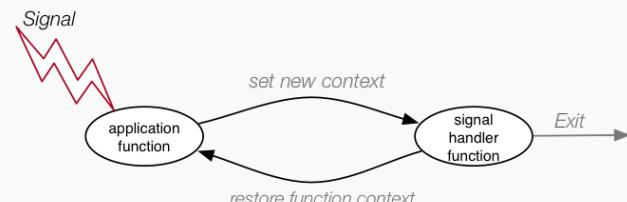
- `SIGILL` ... illegal instruction (**Term** by default)
- `SIGABRT` ... generated by `abort()` (**Core** by default)
- `SIGFPE` ... floating point exception (**Core** by default)
- `SIGSEGV` ... invalid memory reference (**Core** by default)

Signals from external process events, e.g.

- `SIGHUP` ... hangup detected on controlling terminal/process
- `SIGINT` ... interrupt from keyboard (ctrl-c) (**Term** by default)
- `SIGPIPE` ... broken pipe (**Term** by default)
- `SIGCHLD` ... child process stopped or died (**Ign** by default)
- `SIGTSTP` ... stop typed at tty (ctrl-z) (**Stop** by default)

Signal Handlers

Signal Handler = a function invoked in response to a signal - knows which signal it was invoked by - needs to ensure that invoking signal (at least) is blocked - carries out appropriate action; may return



```
#include <signal.h>
typedef void (*sighandler_t)(int);
sighandler_t signal(int signum, sighandler_t handler);
```

- old way to create signal handler - do not use in new code
- set how to handle a signal **signum** (e.g. SIGINT)
- **handler** can be one of ...
 - SIG_IGN ... ignore signal **signum**
 - SIG_DFL ... use default handler for **signum**
 - a user-defined function for **signum** signals
 - function type must be **void (int)**
- returns previous value of signal handler, or SIG_ERR

```
#include <signal.h>
int sigaction(int signum, const struct sigaction *act,
              struct sigaction *oldact);
```

- set how to handle a signal **signum** (e.g. SIGINT)
- **act** defines how signal should be handled
- **oldact** saves a copy of how signal was handled
- if **act->sa_handler == SIG_IGN**, signal is ignored
- if **act->sa_handler == SIG_DFL**, default handler is used
- on success, returns 0; on error, returns -1 and sets **errno**

For much more information: `man 2 sigaction`

Signal Handlers

Details on `struct sigaction` ...

- **void (*sa_handler)(int)**
 - pointer to a handler function, or SIG_IGN or SIG_DFL
- **void (*sa_sigaction)(int, siginfo_t *, void *)**
 - pointer to handler function; used if SA_SIGINFO flag is set
 - allows more context info to be passed to handler
- **sigset_t sa_mask**
 - a mask, where each bit specifies a signal to be blocked
- **int sa_flags**
 - flags to modify how signal is treated
 - (e.g. don't block signal in its own handler)

Signal Handlers

Details on `siginfo_t` ...

- **si_signo** ... signal being handled
- **si_errno** ... any **errno** value associated with signal
- **si_pid** ... process ID of sending process
- **si_uid** ... user ID of owner of sending process
- **si_status** ... exit value for process termination
- etc. etc. etc.

For more details: [bits/types/siginfo_t.h](#) (system-dependent)

```
#include <signal.h>
void signal_handler(int signum) {
    printf("signal number %d received\n", signum);
}

int main(void) {
    struct sigaction action = {.sa_handler = signal_handler};
    sigaction(SIGUSR1, &action, NULL);
    printf("I am process %d waiting for signal %d\n", getpid(), SIGUSR1);
    // loop waiting for signal
    // bad consumes CPU/electricity/battery
    // sleep would be better
    while (1) {
    }
}
```

source code for busy_wait_for_signal.c

9

```
#include <unistd.h>
unsigned int sleep(unsigned int seconds);



- sleep() suspended the caller for seconds of real-time
- efficient way to wait for an event such as an signal
- allows operating system to run other processes

```

10

```
#include <signal.h>
void signal_handler(int signum) {
    printf("signal number %d received\n", signum);
}

int main(void) {
    struct sigaction action = {.sa_handler = signal_handler};
    sigaction(SIGUSR1, &action, NULL);
    printf("I am process %d waiting for signal %d\n", getpid(), SIGUSR1);
    // suspend execution for 1 hour
    sleep(3600);
}
```

source code for wait_for_signal.c

```
include <sys/types.h>
#include <signal.h>
int kill(pid_t pid, int sig);



- send signal number sig to process number pid
- if successful, return 0; on error, return -1 and set errno



int main(int argc, char *argv[]) {
    if (argc != 3) {
        fprintf(stderr, "Usage: %s <signal> <pid>\n", argv[0]);
        return 1;
    }
    int signal = atoi(argv[1]);
    int pid = atoi(argv[2]);
    kill(pid, signal);
}
```

source code for send_signal.c

11

12

example - ignoring a signal

```
#include <signal.h>
int main(void) {
    // catch SIGINT which is sent if user types ctrl-d
    struct sigaction action = {.sa_handler = SIG_IGN};
    sigaction(SIGINT, &action, NULL);
    while (1) {
        printf("Can't interrupt me, I'm ignoring ctrl-C\n");
        sleep(1);
    }
}
```

source code for ignore_control_c.c

example - a simple signal handler

```
#include <signal.h>
void ha_ha(int signum) {
    printf("Ha Ha!\n"); // I/O can be unsafe in a signal handler
}
int main(void) {
    // catch SIGINT which is sent if user types ctrl-d
    struct sigaction action = {.sa_handler = ha_ha};
    sigaction(SIGINT, &action, NULL);
    while (1) {
        printf("Can't interrupt me, I'm ignoring ctrl-C\n");
        sleep(1);
    }
}
```

source code for laugh_at_control_c.c

13

14

example - another simple signal handler

```
#include <signal.h>
int signal_received = 0;
void stop(int signum) {
    signal_received = 1;
}
int main(void) {
    // catch SIGINT which is sent if user types ctrl-C
    struct sigaction action = {.sa_handler = stop};
    sigaction(SIGINT, &action, NULL);
    while (!signal_received) {
        printf("Type ctrl-c to stop me\n");
        sleep(1);
    }
    printf("Good bye\n");
}
```

source code for stop_with_control_c.c

example - catching an internal error with a signal handler

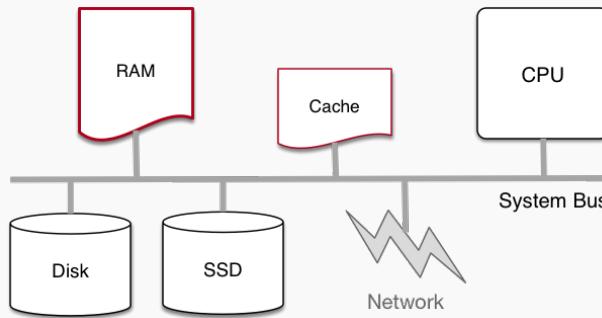
```
#include <signal.h>
#include <stdlib.h>
void report_signal(int signum) {
    printf("Signal %d received\n", signum);
    printf("Please send help\n");
    exit(0);
}
int main(int argc, char *argv[]) {
    struct sigaction action = {.sa_handler = report_signal};
    sigaction(SIGFPE, &action, NULL);
    // this will produce a divide by zero
    // if there are no command-line arguments
    // which will cause program to receive SIGFPE
    printf("%d\n", 42/(argc - 1));
    printf("Good bye\n");
}
```

15

16

source code for catch_error.c

Systems typically contain 4-16GB of volatile RAM



Plus a hierarchy of smaller cache memory -on or off the CPU chip.

- Many small embedded system run without operating system.
- Single program running, probably written in C.
- Devices (sensors, switches, ...) often wired at particular address.
- E.g. can set motor speed by storing byte at 0x100400.
- Program accesses (any) RAM directly.
- Development and debugging tricky.
- Widely used for simple micro-controllers.
- Parallelism and exploiting multiple-core CPUs problematic

1

2

Single Process Resident in RAM with Operating System

- Operating system need (simple) hardware support.
- Part of RAM (kernel space) must be accessible only in a privileged mode.
- System call enables privileged mode and passes execution to operating system code in kernel space.
- Privileged mode disabled when system call returns.
- Privileged mode could be implemented by a bit in a special register
- If only one process resident in RAM at any time - switching between processes is slow.
- Operating system must write out all memory of old process to disk and read all memory of new process from disk.
- OK for some uses, but inefficient in general.
- Little used in modern computing.

Multi Processes Resident in RAM without Virtual Memory

- If multiple processes to be resident in RAM O/S can swap execution between them quickly.
- RAM belonging to other processes & kernel must be protected
- Hardware support can limit process accesses to particular **segment** (region) of RAM.
- BUT program may be loaded anywhere in RAM to run
- Breaks instructions which use absolute addresses, e.g.: `lw, sw, jr`
- Either programs can't use absolute memory addresses (relocatable code)
- Or code has to be modified (relocated) before it is run - not possible for all code!
- Major limitation - much better if programs can assume always have same address space
- Little used in modern computing.

3

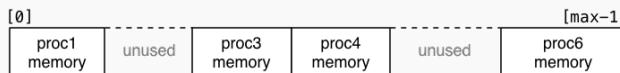
4

Virtual Memory

- Big idea - disconnect address processes use from actual RAM address.
- Operating system translates (virtual) address a process uses to an physical (actual) RAM address.
- Convenient for programming/compilers - each process has same virtual view of RAM.
- Can have multiple processes be in RAM, allowing fast switching
- Can load part of processes into RAM on demand.
- Provides a mechanism to share memory between processes.
- Address to fetch every instruction to be executed must be translated.
- Address for load/store instructions (e.g. `lw, sw`) must be translated .
- Translation needs to be really fast so largely implemented in hardware (silicon).

Virtual Memory with One Memory Segment Per Process

Consider a scenario with multiple processes loaded in memory:

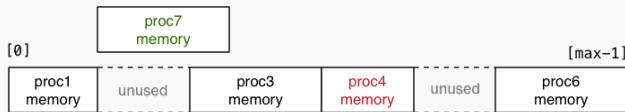


- Every process is in a contiguous section of RAM, starting at address **base** finishing at address **limit**.
- Each process sees its own address space as $[0 .. \text{size} - 1]$
- Process can be loaded anywhere in memory without change.
- Process accessing memory address **a** is translated to $\mathbf{a} + \mathbf{base}$
- and checked that $\mathbf{a} + \mathbf{base}$ is $< \mathbf{limit}$ to ensure process only access its memory
- Easy to implement in hardware.

5

Virtual Memory with One Memory Segment Per Process

Consider the same scenario, but now we want to add a new process



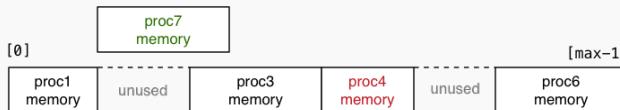
- The new process doesn't fit in any of the unused slots (fragmentation).
- Could move some process to make a single large slot



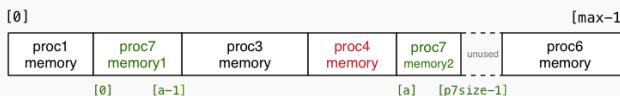
- Slow if RAM heavily used.
- Does not allow sharing or loading on demand.
- Limits process address space to size of RAM.
- Little used in modern computing.

Virtual Memory with Multiple Memory Segments Per Process

Idea: split process memory over multiple parts of physical memory.



becomes



7

8

Virtual Memory with Multiple Memory Segments Per Process

With arbitrary sized memory segments, translating virtual to physical address is complicated making hardware support difficult:

```
// translate virtual_address to physical RAM address
uint32_t translate(uint32_t process_id, uint32_t virtual_addr) {
    uint32_t n_segments;
    Segment *segments = get_segments(process_id, &n_segments);
    for (int i = 0; i < n_segments; i++) {
        Segment *c = &segments[i];
        if (virtual_addr >= c->base &&
            virtual_addr < c->base + c->size) {
            uint32_t offset = virtual_addr - c->base;
            return c->mem + offset;
        }
    }
    // handle illegal memory access
}
```

Virtual Memory with Pages

With pages, translating virtual to physical address is simpler making hardware support difficult:

```
// translate virtual_address to physical RAM address
uint32_t translate(uint32_t process_id, uint32_t virtual_addr) {
    uint32_t pt_size;
    PageInfo *page_table = get_page_table(process_id, &pt_size);
    page_number = virtual_addr / PAGE_SIZE;
    if (page_number < pt_size) {
        uint32_t offset = virtual_addr % PAGE_SIZE;
        return PAGE_SIZE * page_table[page_number].frame + offset;
    }
    // handle illegal memory access
}
```

Virtual Memory with Pages

Address mapping would be simpler if all segments were same size

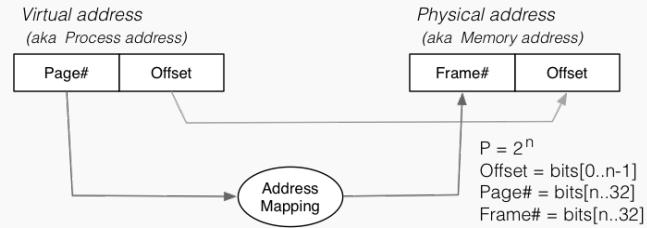
- call each segment of address space a **page**
- make all pages the same size P
- page I holds addresses: $I*P \dots (I+1)*P$
- translation of addresses can be implemented with an array
- each process has an array called the **page table**
- each array element contains the physical address in RAM of that page
- for virtual address V , **page_table[V / P]** contains physical address of page
- the address will be at offset $V \% P$ in both pages
- so physical address for V is: $\text{page_table}[V / P] + V \% P$

9

10

Address Mapping

If $P == 2^n$, then address mapping becomes



- Calculation of *page_number* and *offset* can be faster/simpler bit operations if $PAGE_SIZE == 2^n$, e.g. 4096, 8192, 16384
- Note *PageInfo* entries will have more information about the page ...

11

12

A side-effect of this type of virtual → physical address mapping

- don't need to load all of process's pages up-front
- start with a small memory "footprint" (e.g. `main` + stack top)
- load new process address pages into memory *as needed*
- grow up to the size of the (available) physical memory

The strategy of ...

- dividing process memory space into fixed-size pages
- on-demand loading of process pages into physical memory

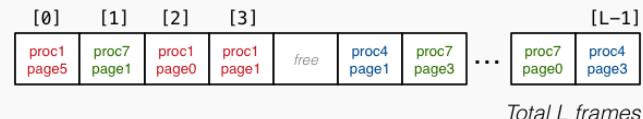
is what is generally meant by *virtual memory*

Pages/frames are typically 4KB .. 256KB in size

With 4GB memory, would have $\approx 1 \text{ million} \times 4\text{KB}$ frames

Each frame can hold one page of process address space

Leads to a memory layout like this (with L total pages of physical memory):



When a process completes, all of its frames are released for re-use

13

14

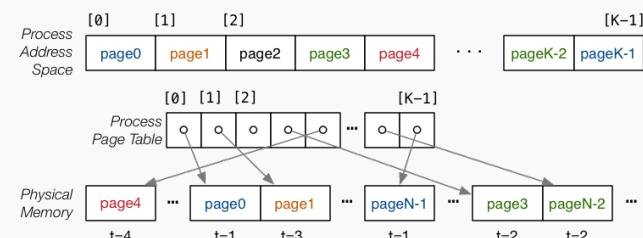
Page Tables

Example Page Table

Consider a possible per-process page table, e.g.

- each page table entry (PTE) might contain
 - page status ... *not_loaded*, *loaded*, *modified*
 - frame number of page (if *loaded*)
 - ... maybe others ... (e.g. last accessed time)
- we need $\lceil \text{ProcSize}/\text{PageSize} \rceil$ entries in this table

Example of page table for one process:



Timestamps show when page was loaded.

15

16

```

typedef struct {int status, int frame, ...} PageInfo;

uint32_t translate(uint32_t process_id, uint32_t virtual_addr) {
    uint32_t pt_size;
    PageInfo *page_table = get_page_table(process_id, &pt_size);
    page_number = virtual_addr / PAGE_SIZE;
    if (page_number < pt_size) {
        if (page_table[page_number].status != LOADED) {
            // page fault - need to load page into free frame
            page_table[page_number].frame = ???;
            page_table[page_number].status = LOADED;
        }
        uint32_t offset = virtual_addr % PAGE_SIZE;
        return PAGE_SIZE * page_table[page_number].frame + offset;
    }
    // handle illegal memory access
}

```

17

Consider a new process commencing execution ...

- initially has zero pages loaded
- load page containing code for `main()`
- load page for `main()`'s stack frame
- load other pages when process references address within page

Do we ever need to load all process pages at once?

18

From observations of running programs ...

- in any given window of time, process typically access only a small subset of their pages
- often called *locality of reference*
- subset of pages called the *working set*

Implications:

- if each process has a relatively small working set,
can hold pages for many active processes in memory at same time
- if only need to hold some of process's pages in memory,
process address space can be larger than physical memory

We say that we "load" pages into physical memory

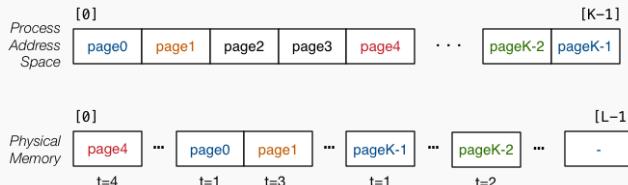
But where are they loaded from?

- code is loaded from the executable file stored on disk into read-only pages
- some data (e.g. C strings) also loaded into read-only pages
- initialised data (C global/static variables) also loaded from executable file
- pages for uninitialised data (heap, stack) are zero-ed
 - prevents information leaking from other processes
 - results in uninitialised local (stack) variables often containing 0

Consider a process whose address space exceeds physical memory

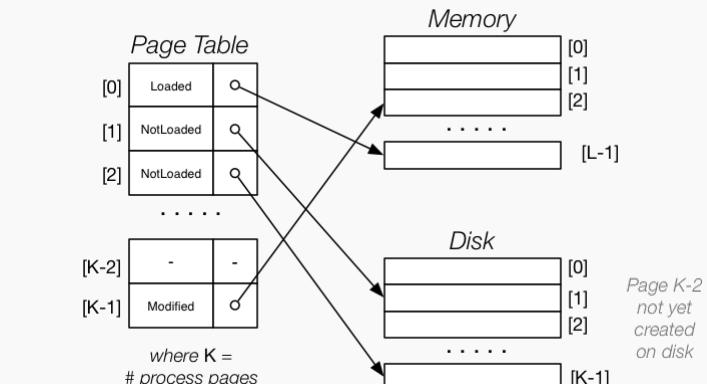
We can imagine that a process's address space ...

- exists on disk for the duration of the process's execution
- and only some parts of it are in memory at any given time



Transferring pages between disk↔memory is **very** expensive

- need to ensure minimal reading from / writing to disk



21

22

Virtual Memory - Handling Page Faults

An access to a page which is not-loaded in RAM is called a **page fault**.

Where do we load it in RAM?

First need to check for a free frame

- need a way of quickly identifying free frames
- commonly handled via a free list

What if there are currently no free page frames, possibilities:

- suspend the requesting process until a page is freed
- replace one of the currently loaded/used pages

Suspending requires the operating system to

- mark the process as unable to run until page available
- switch to running another process
- mark the process as able to run when page available

Page Replacement

If no free pages we need to choose a page to evict:

- best page is one that won't be used again by its process
- prefer pages that are read-only (no need to write to disk)
- prefer pages that are unmodified (no need to write to disk)
- prefer pages that are used by only one process (see later)

OS can't predict whether a page will be required again by its process

But we do know whether it has been used recently (if we record this)

One good heuristic - replace Least Recently Used (LRU) page.

- page not used recently probably not needed again soon

23

24

Show how the page frames and page tables change when

- there are 4 page frames in memory
- the process has 6 pages in its virtual address space
- a LRU page replacement strategy is used

For each of the following sequences of virtual page accesses

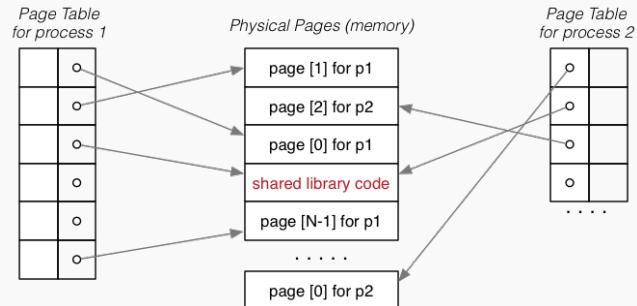
0, 5, 0, 0, 5, 1, 5, 1, 2, 4, 3, 3, 4, 2, 5, 3, 2

5, 0, 0, 0, 5, 1, 1, 5, 1, 5, 2, 2, 3, 0, 0, 5

Assume that all PTEs and frames are initially empty/unused

Virtual memory allows sharing of read-only pages (e.g. library code)

- several processes include same frame in virtual address space

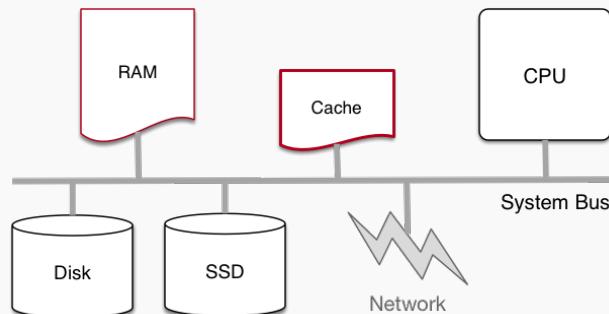


25

26

Cache Memory

Cache memory = small, fast memory* close to CPU*



Small = MB, Fast = 5 × RAM

Cache Memory

Cache memory

- holds parts of RAM that are (hopefully) heavily used
- transfers data to/from RAM in blocks (*cache blocks*)
- memory reference hardware first looks in cache
 - if required address is there, use its contents
 - if not, get it from RAM and put in cache
 - possibly replacing an existing cache block
- replacement strategies have similar issues to virtual memory

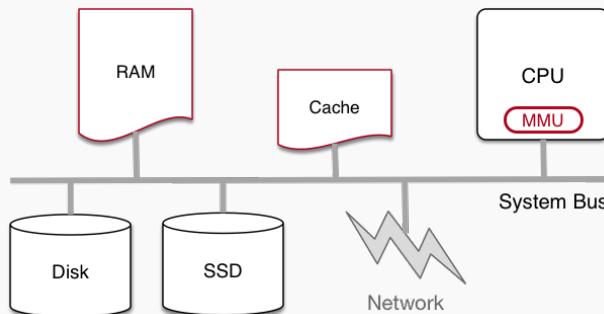
27

28

Memory Management Hardware

Address translation is very important/frequent

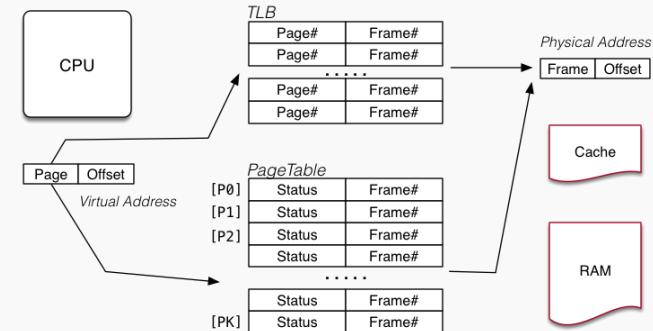
- provide specialised hardware (MMU) to do it efficiently
- sometimes located on CPU chip, sometimes separate



Memory Management Hardware

TLB = translation lookaside buffer

- lookup table containing (virtual,physical) address pairs



Concurrency multiple computations in overlapping time periods;
does not have to be simultaneous

Parallelism multiple computations executing simultaneously

Parallel computation occurs at different level:

- spread across computers (e.g., with MapReduce)
- multiple cores of a CPU executing different instructions (MIMD)
- multiple cores of a CPU executing same instruction (SIMD)
 - e.g. GPU rendering pixels

Both parallelism and concurrency need to deal with *synchronisation*.

Example: *Map-reduce* is a popular programming model for

- manipulating very large data sets
- on a large network of computers (local or distributed)

The *map* step filters data and distributes it to nodes

- data distributed as $(key, value)$ pairs
- each node receives a set of pairs with common *key(s)*

Nodes then perform calculation on received data items

The *reduce* step computes the final result

- by combining outputs (calculation results) from the nodes

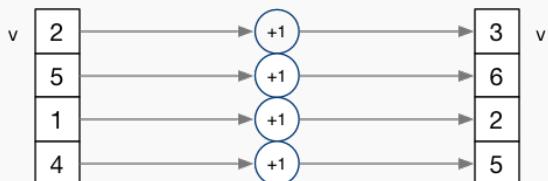
Also needs a way to determine when all calculations completed

1

2

Parallelism Across a an Array

- multiple identical processors
- each given one element of a data structure from main memory
- each performing same computation on that element (SIMD)
- results copied back to data structure in main memory



But not totally independent: need to *synchronise* on completion

Example: GPU rendering pixels or neural network

Parallelism Across Processes

One method for creating parallelism:

Use `posix_spawn()` to create multiple processes, each does part of job.

- child executes concurrently with parent
- runs in its own address space
- inherits some state information from parent, e.g. open fd's

Processes have some disadvantages

- process switching expensive
- each require a significant amount of state (RAM)
- communication between processes limited and/or slow

One big advantage - separate address spaces make processes more robust.

3

4

Parallelism within Processes

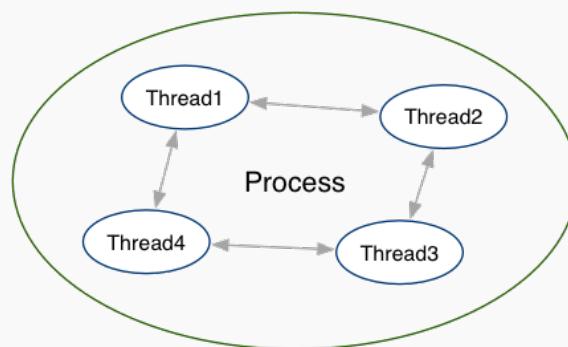
threads - mechanism for parallelism within process.

- threads allow simultaneous execution within process
- each thread has its own execution state
- threads within a process have same address space:
 - threads share code (functions)
 - threads share global & static variables
 - threads share heap (`malloc`)
- but separate stack for each thread
 - local variables not shared
- threads share file descriptor
- threads share signals

POSIX threads (pThreads)

// POSIX threads widely supported in Unix-like
// and other systems (Windows). Provides functions
// to create/synchronize/destroy... threads

```
#include <pthread.h>
```



5

6

Create A POSIX Thread

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine)(void *),  
                  void *arg);
```

- creates a new thread with attributes specied in `attr`
 - `attr` can be NULL
- thread info stored in `*thread`
- thread starts by executing `start_routine(arg)`
- returns 0 if OK, -1 otherwise and sets `errno`
- analogous to `posix_spawn()`

Wait for A POSIX Thread

```
int pthread_join(pthread_t thread, void **retval)
```

- wait until `thread` terminates
- `thread` return (or `pthread_exit()`) value is placed in `*retval`
- if `thread` has already exited, does not wait
- if main returns or exit called, all threads terminated
- programs typically need to wait for all threads before main returns/exit called
- analogous to `waitpid`

7

8

- ```
void pthread_exit(void *retval);;
```
- terminate execution of thread (and free resources)
  - **retval** is returned (see `pthread_join`)
  - if **thread** has already exited, does not wait
  - analogous to `exit`

```
#include <pthread.h>
// this function is called to start thread execution
// it can be given any pointer as argument (int *) in this example
void *run_thread(void *argument) {
 int *p = argument;
 for (int i = 0; i < 10; i++) {
 printf("Hello this is thread #%d: i=%d\n", *p, i);
 }
 // a thread finishes when the function returns or thread_exit
 // a pointer of any type can be returned
 // this can be obtained via thread_join's 2nd argument
 return NULL;
}
```

source code for two\_threads.c

9

10

```
//create two threads performing almost the same task
pthread_t thread_id1;
int thread_number1 = 1;
pthread_create(&thread_id1, NULL, run_thread, &thread_number1);
int thread_number2 = 2;
pthread_t thread_id2;
pthread_create(&thread_id2, NULL, run_thread, &thread_number2);
// wait for the 2 threads to finish
pthread_join(thread_id1, NULL);
pthread_join(thread_id2, NULL);
```

source code for two\_threads.c

```
pthread_t thread_id1;
int thread_number = 1;
pthread_create(&thread_id1, NULL, run_thread, &thread_number);
thread_number = 2;
pthread_t thread_id2;
pthread_create(&thread_id2, NULL, run_thread, &thread_number);
pthread_join(thread_id1, NULL);
pthread_join(thread_id2, NULL);
```

source code for two\_threads\_broken.c

- variable **thread\_number** will probably have changed in main before thread 1 starts executing
- so thread 1 will probably print `Hello this is thread 2`

## Simple example - Creating Many threads

```
int n_threads = strtol(argv[1], NULL, 0);
assert(n_threads > 0 && n_threads < 100);
pthread_t thread_id[n_threads];
int argument[n_threads];
for (int i = 0; i < n_threads; i++) {
 argument[i] = i;
 pthread_create(&thread_id[i], NULL, run_thread, &argument[i]);
}
// wait for the threads to finish
for (int i = 0; i < n_threads; i++) {
 pthread_join(thread_id[i], NULL);
}
```

source code for n\_threads.c

## Simple example - Dividing a task between threads

```
struct job {
 long start;
 long finish;
 double sum;
};

void *run_thread(void *argument) {
 struct job *j = argument;
 long start = j->start;
 long finish = j->finish;
 double sum = 0;
 for (long i = start; i < finish; i++) {
 sum += i;
 }
 j->sum = sum;
}
```

source code for thread\_sum.c

13

14

## Simple example - Dividing a task between threads

```
printf("Creating %d threads to sum the first %lu integers\n",
 n_threads, integers_to_sum);
printf("Each thread will sum %lu integers\n", integers_per_thread);
pthread_t thread_id[n_threads];
struct job jobs[n_threads];
for (int i = 0; i < n_threads; i++) {
 jobs[i].start = i * integers_per_thread;
 jobs[i].finish = jobs[i].start + integers_per_thread;
 if (jobs[i].finish > integers_to_sum) {
 jobs[i].finish = integers_to_sum;
 }
 // create a thread which will sum integers_per_thread integers
 pthread_create(&thread_id[i], NULL, run_thread, &jobs[i]);
}
```

source code for thread\_sum.c

## Simple example - Dividing a task between threads

```
double overall_sum = 0;
for (int i = 0; i < n_threads; i++) {
 pthread_join(thread_id[i], NULL);
 overall_sum += jobs[i].sum;
}
//
printf("\nCombined sum of integers 0 to %lu is %.0f\n",
 integers_to_sum, overall_sum);
```

source code for thread\_sum.c

15

16

```

double overall_sum = 0;
for (int i = 0; i < n_threads;i++) {
 pthread_join(thread_id[i], NULL);
 overall_sum += jobs[i].sum;
}
//
printf("\nCombined sum of integers 0 to %lu is %.0f\n",
 integers_to_sum, overall_sum);

```

source code for thread\_sum.c

- on a AMD Ryzen 3900x which can run 24 threads simultaneously
- 1 thread takes 6.9 seconds to sum first 10000000000 integers
- 2 threads takes 3.6 seconds to sum first 10000000000 integers
- 4 threads takes 1.8 seconds to sum first 10000000000 integers
- 12 threads takes 0.6 seconds to sum first 10000000000 integers
- 24 threads takes 0.3 seconds to sum first 10000000000 integers
- 50 threads takes 0.3 seconds to sum first 10000000000 integers
- 500 threads takes 0.3 seconds to sum first 10000000000 integers

17

18

## Example - Unsafe Access to Global Variable

```

int bank_account = 0;
// add $1 to Andrew's bank account 100,000 times
void *add_100000(void *argument) {
 for (int i = 0; i < 100000; i++) {
 // execution may switch threads in middle of assignment
 // between load of variable value
 // and store of new variable value
 // changes other thread makes to variable will be lost
 nanosleep(&(struct timespec){.tv_nsec = 1}, NULL);
 bank_account = bank_account + 1;
 }
 return NULL;
}

```

source code for bank\_account\_broken.c

## Example - Unsafe Access to Global Variable

```

int main(void) {
 //create two threads performing the same task
 pthread_t thread_id1;
 pthread_create(&thread_id1, NULL, add_100000, NULL);
 pthread_t thread_id2;
 pthread_create(&thread_id2, NULL, add_100000, NULL);
 // wait for the 2 threads to finish
 pthread_join(thread_id1, NULL);
 pthread_join(thread_id2, NULL);
 // will probably be much less than $200000
 printf("Andrew's bank account has %d\n", bank_account);
 return 0;
}

```

source code for bank\_account\_broken.c

19

20

## Global Variable: Race Condition

Incrementing a global variable is not an atomic (indivisible) operation.

```
int bank_account;
void *thread(void *a) {
 // ...
 bank_account++;
 // ...
}
```

.data  
bank\_account: .word 0

la \$t0, bank\_account  
lw \$t1, (\$t0)  
# \$t1 == 42  
addi \$t1, \$t1, 1  
# \$t1 == 43  
sw \$t1, (\$t0)  
# bank\_account == 43

## Global Variable: Race Condition

If bank\_account == 42 and two threads increment simultaneously.

la \$t0, bank\_account  
lw \$t1, (\$t0)  
# \$t1 == 42  
addi \$t1, \$t1, 1  
# \$t1 == 43  
sw \$t1, (\$t0)  
# bank\_account == 43

la \$t0, bank\_account  
lw \$t1, (\$t0)  
# \$t1 == 42  
addi \$t1, \$t1, 1  
# \$t1 == 43  
sw \$t1, (\$t0)  
# bank\_account == 43

One increment is lost.

Note threads don't share registers or stack (local variable).

They do share global variables.

21

22

## Global Variable: Race Condition

If bank\_account == 100 and two threads change it simultaneously.

```
la $t0, bank_account
lw $t1, ($t0)
$t1 == 100
addi $t1, $t1, 100
$t1 == 200
sw $t1, ($t0)
bank_account == ?
```

la \$t0, bank\_account  
lw \$t1, (\$t0)  
# \$t1 == 100  
addi \$t1, \$t1, -50  
# \$t1 == 50  
sw \$t1, (\$t0)  
# bank\_account == 50 or 200

## Exclude Other Threads from Code

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- only one thread can enter a *critical section*
- establishes mutual exclusion — mutex
- call `pthread_mutex_lock` before
- call `pthread_mutex_unlock` after
- only 1 thread can execute in protected code
- for example:

```
pthread_mutex_lock(&bank_account_lock);
andrews_bank_account += 1000000;
pthread_mutex_unlock(&bank_account_lock);
```

23

24

## Example - Protecting Access to Global Variable with a Mutex

```
int bank_account = 0;
pthread_mutex_t bank_account_lock = PTHREAD_MUTEX_INITIALIZER;
// add $1 to Andrew's bank account 100,000 times
void *add_100000(void *argument) {
 for (int i = 0; i < 100000; i++) {
 pthread_mutex_lock(&bank_account_lock);
 // only one thread can execute this section of code at any
 bank_account = bank_account + 1;
 pthread_mutex_unlock(&bank_account_lock);
 }
 return NULL;
}
```

source code for bank\_account\_mutex.c

## Semaphores

Semaphores are special variables which provide a more general synchronisation mechanism than mutexes.

```
#include <semaphore.h>
int sem_init(sem_t *sem, int pshared,
 unsigned int value);
int sem_post(sem_t *sem);
int sem_wait(sem_t *sem);

■ sem_init initialises sem to value
■ sem_wait - classically called P()
 ■ if sem > 0, decrement sem and continue
 ■ otherwise, wait until sem > 0
■ sem_post - classically called V()
 ■ increment sem and continue
```

25

26

## Allow n threads access to a resource

```
#include <semaphore.h>
sem_t sem;
sem_init(&sem, 0, n);

sem_wait(&sem);
// only n threads can be in executing
// in here simultaneously
sem_post(&sem);
```

## Protecting Access to Global Variable with a Semaphore

```
sem_t bank_account_semaphore;
// add $1 to Andrew's bank account 100,000 times
void *add_100000(void *argument) {
 for (int i = 0; i < 100000; i++) {
 // decrement bank_account_semaphore if > 0
 // otherwise wait until > 0
 sem_wait(&bank_account_semaphore);
 // only one thread can execute this section of code at any
 // because bank_account_semaphore was initialized to 1
 bank_account = bank_account + 1;
 // increment bank_account_semaphore
 sem_post(&bank_account_semaphore);
 }
 return NULL;
}
```

27

28

source code for bank\_account\_semaphore.c

## Protecting Access to Global Variable with a Semaphore

```
// initialize bank_account_semaphore to 1
sem_init(&bank_account_semaphore, 0, 1);
//create two threads performing the same task
pthread_t thread_id1;
pthread_create(&thread_id1, NULL, add_100000, NULL);
pthread_t thread_id2;
pthread_create(&thread_id2, NULL, add_100000, NULL);
// wait for the 2 threads to finish
pthread_join(thread_id1, NULL);
pthread_join(thread_id2, NULL);
// will always be $200000
printf("Andrew's bank account has %d\n", bank_account);
sem_destroy(&bank_account_semaphore);
```

source code for bank\_account\_semaphore.c

## File Locking

If a process tries to acquire a *shared lock* ...

- if file not locked or other shared locks, OK
- if file has exclusive lock, blocked

If a process tries to acquire an *exclusive lock* ...

- if file is not locked, OK
- if any locks (shared or exclusive) on file, blocked

If using a non-blocking lock

- flock() returns 0 if lock was acquired
- flock() returns -1 if process would have been blocked

## File Locking

```
int flock(int FileDesc, int Operation)
```

Similar to mutexes for a file.

- controls access to shared files (**note:** files not fds)
- possible operations
  - LOCK\_SH ... acquire shared lock
  - LOCK\_EX ... acquire exclusive lock
  - LOCK\_UN ... unlock
  - LOCK\_NB ... operation fails rather than blocking
- in blocking mode, flock() does not return until lock available
- only works correctly if all processes accessing file use locks
- return value: 0 in success, -1 on failure

29

30

## Concurrent Programming is Complex

Concurrency is *complex* with many issues beyond this course:

**Data races** thread behaviour depends on unpredictable ordering;  
can produce difficult bugs or security vulnerabilities

**Deadlock** threads stopped because they are wait on each other

**Livelock** threads running without making progress

**Starvation** threads never getting to run

31

32

## Example - deadlock accessing two resources

```
void *swap1(void *argument) {
 for (int i = 0; i < 100000; i++) {
 pthread_mutex_lock(&bank_account1_lock);
 pthread_mutex_lock(&bank_account2_lock);
 int tmp = andrews_bank_account1;
 andrews_bank_account1 = andrews_bank_account2;
 andrews_bank_account2 = tmp;
 pthread_mutex_unlock(&bank_account2_lock);
 pthread_mutex_unlock(&bank_account1_lock);
 }
 return NULL;
}
```

source code for bank\_account\_deadlock.c

## Example - deadlock accessing two resources

```
void *swap2(void *argument) {
 for (int i = 0; i < 100000; i++) {
 pthread_mutex_lock(&bank_account2_lock);
 pthread_mutex_lock(&bank_account1_lock);
 int tmp = andrews_bank_account1;
 andrews_bank_account1 = andrews_bank_account2;
 andrews_bank_account2 = tmp;
 pthread_mutex_unlock(&bank_account1_lock);
 pthread_mutex_unlock(&bank_account2_lock);
 }
 return NULL;
}
```

source code for bank\_account\_deadlock.c

33

34

## Example - deadlock accessing two resources

```
int main(void) {
 //create two threads performing almost the same task
 pthread_t thread_id1;
 pthread_create(&thread_id1, NULL, swap1, NULL);
 pthread_t thread_id2;
 pthread_create(&thread_id2, NULL, swap2, NULL);
 // threads will probably never finish
 // deadlock will likely likely occur
 // with one thread holding bank_account1_lock
 // and waiting for bank_account2_lock
 // and the other thread holding bank_account2_lock
 // and waiting for bank_account1_lock
 pthread_join(thread_id1, NULL);
 pthread_join(thread_id2, NULL);
 return 0;
}
```

source code for bank\_account\_deadlock.c

35

## Course Goals

At the end of COMP1521, we hope that you ...

- understand the structure of computer systems
- can describe how computers/programs work at a low-level
- are better able to reason about and debug your C programs

Major topics ...

- components of modern computer systems
- how C programs execute (at the machine level)
- how to write (MIPS) assembly language
- Unix/Linux system-level programming

## Syllabus/Topics

- bit-level operations
- representation of integers & doubles
- the basic components of a (MIPS) CPU
- representation of programs as (MIPS) machine code
- how to write programs in (MIPS) assembler
- how C programs are implemented as (MIPS) instructions
- systems programming, including:
  - file operations
  - processes
- representation of characters as Unicode
- introduction to virtual memory
- introduction to threads/concurrency

1

2

## Assessment

## Labs and Tests

- 15% Labs
- 10% Weekly Programming Tests
- 15% Assignment 1 — due week 7
- 15% Assignment 2 — due week 10
- 45% Final Exam

Above marks may be scaled to ensure an appropriate distribution

**To pass you must:**

- score 50/100 overall
- score 18/45 on final exam

For example:

55/100 overall, 17/45 on final exam ⇒ **55 UF** not 55 PS

- 9 labs - weeks 1-5,7-10
  - max lab mark 2 marks with challenge exercises
  - max lab mark ~1.6 marks without challenge exercises
  - 9 labs marks summed and capped at 15
  - you could get 99% for lab mark without challenge exercises
  - most people will get 12+/15
- 8 tests - weeks 3...10:
  - max test mark 1.7
  - best 6 of 8 test marks summed and capped at to give mark out of 10.
  - most people will get 7+/10

3

4

- The 20T2 COMP1521 Final Exam is available at:  
<https://cgi.cse.unsw.edu.au/~cs1521/20T3/exam/practice/questions>
- You can complete it as a practice exam
- Autotests available
- Sample answers will be released noon Wednesday 25th
- 20T3 exam will use similar format similar for at least some questions

- Run under same conditions as Weekly Tests
- Except 3 hours and some questions may not be coding
- Saturday 28 November 13:00 — 16:00
- Exam will be released on class web site at 12:50
- Announcements before & during exam will be sent to your UNSW email
- Questions during exam can be sent to [cs1521.exam@cse.unsw.edu.au](mailto:cs1521.exam@cse.unsw.edu.au)
- We may also announce a 2nd place you can ask questions during exam.

5

6

## Exam Conditions

## Exam Format

- You are not permitted to communicate (email, phone, message, talk, ...) to anyone but COMP1521 staff during exam
- You are not permitted to get help from anyone but COMP1521 staff during the exam.
- This is a closed book exam: you are not permitted to access papers, books, files on your computer or the internet
- You are permitted to access the exam web pages on the class web site
- You are permitted to access the online language cheatsheets & documentation on the class web site
- Deliberate violation of exam conditions will be treated as serious misconduct

- 10-15 questions
- Each question answered in a separate file.
- Some questions will ask you to write C.
- Some questions will ask you to write MIPS.
- Other languages not permitted (e.g., Python, C++, Java, Rust, ...)
- Some questions may not involve coding
- Answers will be submitted with give.
- Questions not equal difficulty
- Questions may not be worth equal marks

7

8

- Answers must be an specified file, e.g. q1.txt
- Question may specify format of file
  - e.g. 5 integers one per line
  - follow this format **EXACTLY**
- Question must give you an initial file to complete
- File will be submitted with give.

For question that require you to write C or MIPS ...

- Questions will usually include examples.
- You may or may not be given starting code.
- You may or may not be given test data or other files
- 1 or more autotests may be available on submission.
- **Passing autotests does not guarantee any marks.**  
Do your own testing.
- There may be no submission tests for some questions.
- It is not sufficient to match any supplied examples.

9

10

## Marking of Coding Questions

## Special Exam Conditions

- Answers will be run through automatic marking software.
  - Please follow the input/output format shown exactly.
  - Please make your program behave exactly as specified.
- All answers are hand marked, guided by automarking.
  - No marks awarded for style or comments ...
  - But use decent formatting so the marker can read the program!
  - Comments only necessary to tell the marker something.
- Minor errors will result in only a small penalty.
  - e.g. an answer correct except for a missing semi-colon would receive almost full marks.
- No marks will given unless an answer contains a substantial part of a solution (> 33%).
- No marks just for starting a question and writing some code

- any extra time specified in your ELS exam conditions is allowed in this exam
- all student see the same exam questions text
- the text shows the standard exam deadline, any extra time is additional to it
- give should be configured to allow your extra time
- give should show a deadline **including** your extra time
- email `cs1521@cse.unsw.edu.au` immediately during exam if you have concerns regarding ELS conditions
- if ELS conditions prevent you taking exam, alternative will be supp in January

11

12

- a small number of students have exams on the same day
  - MARK1012 & COMP9517
- these exams are timetabled as all day exams
- special considerations have advised they will not give special consideration for clashes with all day exams

- if a problem occurs during the exam, e.g. internet failure
- please document the problem as possible, e.g. take screenshot
- email [cs1521.exam@cse.unsw.edu.au](mailto:cs1521.exam@cse.unsw.edu.au)
- if the problem is of short duration we may be able to give you extra time
- otherwise you'll need to apply for special consideration

13

## Special Consideration — 'Fit to Sit'

## What to study

- By starting the exam, you are saying "I am well enough to sit it".
- If unwell before exam, see a doctor, apply for Special Consideration.
- If you become unwell during the exam
  - email [cs1521.exam@cse.unsw.edu.au](mailto:cs1521.exam@cse.unsw.edu.au)
  - if you can't continue the exam you will need to see a doctor and apply for Special Consideration.

- Important Areas to Focus Your Study On
  - anything covered in a standard lab exercise
  - anything covered in a weekly test
  - anything covered by the assignments
- Less Important Areas
  - challenge lab exercises
  - topics not covered in labs, tests or assignments
  - may still be questions on these topics but not many
- Even Less Important Areas
  - creating pipes and other complex aspects of creating processes
  - complex signal handling
  - semaphores & file locking
  - might or might not be a question on these topics

15

14

16

- Exam marks will be made available via class database when marking is complete.
- I'll send email announcing this.
- Marking will probably take 12 days.
- You will receive marks for individual exam questions.
- You will have opportunity to have marking reviewed.
- Final results will appear on myUNSW.

- UNSW supplementary exams are run centrally.
- Supplementary exams are for students who miss original exam due to illness/misadventure.
- If this is you — apply for special consideration.
- Lecturers & schools can not offer supps.
- Students with borderline results are not offered supps.  
... except potential graduands
- Supp exams centrally timetabled for January 11-15, 2021
- Similar format to final exam

17

## What did you like?

One aim of COMP1521 is to give a taste of many topics:

- Liked MIPS/assembly?  
⇒ COMP3222, COMP3211 ...
- Curious about programming languages?  
⇒ COMP3131, COMP3141, COMP3161, ...
- Liked Operating Systems?  
⇒ COMP3231/3891, COMP9242, ...
- Liked Concurrency?  
⇒ COMP3151, ...
- Liked Networking?  
⇒ COMP3331, COMP4336, COMP4337, ...
- Liked Unix shell?  
⇒ COMP2041

## Course Offerings

COMP1531 Software Engineering Fundamentals

- 2021: term1, term3

COMP2511 Object-oriented Programming

- 2021: term2, term3

COMP2521 Data Structures and Algorithms

- 2021: term1, term2, term3

COMP3231 Operating Systems

- 2021: term1

19

18

20

- MIPS and its relation to C needs to be better explained
- Not enough time to cover (so) many things
- Tuts and Labs need to integrate better.
- Labs a lot of work - but you learnt a lot
- Assignments a **lot** of work - but you learnt a lot

- Most labs exercises (do you agree??)
- Weekly Tests (do you agree??)
- Tutors
- Discourse
- Students

21

22

And that's all ...

- Good Luck
- I hope what you've learnt in this course will be useful.
- I hope you get the mark you deserve.

23